

Causal normalizing flows: from theory to practice

Adrian Javalov Pablo Sanchez-Martín Isabel Valera

Department of Computer Science, Saarland University, Saarbrücken, Germany
Max Planck Institute for Intelligent Systems, Tübingen, Germany
Max Planck Institute for Software Systems, Saarbrücken, Germany

Why We Bridge Causal ML and Generative Modeling

- We study what-if and counterfactual reasoning that drives decisions in healthcare, policy, and fairness.
- We leverage normalizing flows for exact likelihoods and tractable inverses on complex data.
- We aim to make causal inference practical with modern deep generative modeling.

The Problem: How We Learn Causal Models and Interventions from Observations

- We seek identifiable causal models from observational data under minimal assumptions.
- We avoid brittle per-node models and heuristic graph surgery for interventions.
- We handle mixed data and partial graphs while remaining scalable and accurate.

Our Core Idea and Contributions (High-Level)

- We unify SCMs and autoregressive NFs via triangular monotone increasing maps and prove identifiability.
- We design causally consistent NF architectures and show a single-layer abductive ANF works best.
- We implement a general do-operator by manipulating exogenous variables for interventions and counterfactuals.
- We extend to mixed discrete–continuous data and partial graphs (SCC-based), and validate extensively.


Takeaway: We learn an identifiable single-flow model aligned with the causal graph that supports sound, efficient interventions and counterfactuals.

From SCMs to Triangular Flows: Why Structure Matters

- We exploit SCM-induced causal orderings to obtain triangular maps from exogenous u to observed x .
- We align triangular, monotone structure with autoregressive NFs for exact likelihoods and tractable inverses.
- We control Jacobian sparsity via the graph, linking flow dependencies to causal edges.

Key pieces

- u are independent exogenous variables.
- $f : u \rightarrow x$ is a triangular flow aligned with the causal order.
- Masks $M(A)$ enforce parent-only dependencies.
- Jacobian sparsity links the flow to the causal graph.



How We Enforce Causal Graphs in Flows

- We use a topological order π and adjacency A to define valid parent dependencies.
- We enforce masks in the Jacobian to avoid spurious paths inconsistent with the graph.
- We ensure causal consistency at the optimum by respecting these structural constraints.




Figure: Causal graph with an ordering π and adjacency A that guide masking in the flow.

Identifiability: How We Isolate Exogenous Variables with Flows

- If an ANF matches the observational distribution, it recovers exogenous variables up to component-wise invertible transforms.
- At optimum, we obtain causal consistency: Jacobian sparsity matches the true causal graph.
- This connects nonlinear ICA insights to causal flow learning.

Key components and roles


- $T_\theta(x)$ (abductive map): maps $x \rightarrow u$ with triangular Jacobian masked by $(I - A)$.
- f (generative map): maps $u \rightarrow x$; triangular structure enforces parent-only dependencies.
- P_u (base): fully factorized exogenous distribution; independence drives identifiability.
- Causal consistency: at optimum, the Jacobian sparsity of T_θ matches the true graph $(I - A)$.

Design Choices: Generative ($u \rightarrow x$) vs Abductive ($x \rightarrow u$)

- We find generative $u \rightarrow x$ requires multiple masked layers; $L \geq \text{graph diameter}$ to avoid shortcuts.
- We use abductive $x \rightarrow u$ with $L = 1$ and graph masking, which is causally consistent and easy to optimize.
- With only ordering known, we learn sparsity via Jacobian regularization on forbidden entries.

What the masks do

- Graph mask $M(A)$: allows edges only from parents; prevents shortcuts.
- Ordering-only case: we regularize forbidden Jacobian entries to learn sparsity.
- Abductive $x \rightarrow u$ with $L = 1 + M(A)$: causally consistent by construction.



- 1) Abduction: map x to $u = T_\theta(x)$.
- 2) Intervention: set the target $x_i = \alpha$ and adjust the corresponding exogenous u_i so that $f_i(u) = \alpha$ (Dirac constraint on u_i).
- 3) Prediction: map the modified u back through f ($u \rightarrow x$) to sample from $p(x \mid \text{do}(x_i = \alpha))$.
- Note: We operate on P_u (not on structural equations) and avoid heuristic graph surgery.

- We observe correct path breaking after intervention, as predicted by the graph.
- We see ancestor-descendant dependencies adjust consistently with theory.
- We confirm principled interventional behavior of causal NFs.

Setup: SIMPSON NLIN; $\text{do}(x_3 = -1.09)$.




Figure: Observed vs. interventional samples in a Simpson's paradox setup; intervention severs non-causal dependencies.

- We use principled dequantization (add noise $\epsilon \in [0, 1]$) to embed discrete variables into flows.
- We rely on recoverability guarantees to preserve the original discrete distribution.
- We enable causal NFs on real-world datasets with categorical features.

- We group nodes with unknown directions into strongly connected components (SCCs).
- We model intra-SCC with a general NF or ANF and preserve the inter-block DAG structure.
- We obtain block-wise identifiability and causal consistency at the SCC level.




Figure: SCC reduction for partial graphs: intra-SCC modeling with preserved inter-SCC DAG.

- We note uncertain directions (e.g., $x^2 \leftrightarrow x^3$) can imply different causal stories, including confounding.
- We use SCC grouping to avoid premature commitment to a single direction.
- We support robust learning while respecting known structure.




Figure: Ambiguity in partial graphs motivates SCC blocks when directions are uncertain.

- Datasets: synthetic SCMs (chains, fork, collider, triangle, Simpson, LARGEBD) and German Credit.
- Metrics: exact likelihood (KL), ATE RMSE, counterfactual RMSE; runtime for training/eval/sampling.
- Baselines: CAREFL (Author–Year) and VACA (Author–Year).




Figure: Causal graphs for the SCM benchmarks used in our experiments.

- We use abductive $x \rightarrow u$ with $L = 1$ and known A for causal consistency and easy optimization.
- Generative $u \rightarrow x$ respects A but needs depth; sampling is faster, evaluation is slower.
- Abductive trains efficiently (evaluation) and remains accurate for causal queries.

Table: Summary of design choices, induced properties, and time complexity for density evaluation and sampling.

Design Choices	Assumption	Model Properties		Causal Consistency	Time Complexity	
		$u \rightarrow x$	$x \rightarrow u$		Sampling	Evaluation
Generative	Ordering	×	×		$O(L)$	$O(dL)$
$u \rightarrow x$ Generative	Graph A	✓	×		$O(L)$	$O(dL)$
Abductive ($L > 1$)	Ordering	×	×		$O(dL)$	$O(L)$
Abductive ($L > 1$) $x \rightarrow u$	Graph A	×	×		$O(dL)$	$O(L)$
Abductive ($L = 1$)	Graph A	✓	✓	✓	$O(dL)$	$O(L)$

- We observe generative $u \rightarrow x$ needs $L \geq \text{diameter}$; abductive $x \rightarrow u$ works well with $L = 1$.
- We find graph masks significantly improve interventional and counterfactual accuracy.
- We use Jacobian regularization when only ordering is known; it is unnecessary with full masks.




Figure: Ablation on a 4-node chain: effect of flow depth for generative vs. abductive models.

Benchmark Comparisons: Accuracy and Speed

- We match or improve KL vs. CAREFL (Author-Year) and strongly improve ATE/CFRMSE; VACA (Author-Year) is slower and less accurate.
- We observe training/evaluation times favor causal NFs across datasets.

Table: Comparison, on different SCMs, of the proposed Causal NF, VACA, and CAREFL with the do-operator.

Dataset	Model	KL	ATERMSE	CFRMSE	Training	Evaluation	Sampling
3-CHAIN LIN	Causal NF	0.000.00	0.050.01	0.040.01	0.410.06	0.480.10	0.760.06
	CAREFL	0.000.00	0.200.13	0.200.09	0.680.24	0.970.33	1.940.77
	VACA	4.441.03	5.760.07	4.980.10	36.191.54	28.330.72	75.344.58
3-CHAIN NLIN	Causal NF	0.000.00	0.030.01	0.020.01	0.520.06	0.560.03	1.020.05
	CAREFL	0.000.00	0.050.02	0.040.02	0.600.22	0.840.22	1.660.41
	VACA	12.821.00	1.540.03	1.320.02	39.454.12	30.932.30	84.369.60
4-CHAIN LIN	Causal NF	0.000.00	0.070.02	0.040.01	0.560.08	0.620.15	1.540.40
	CAREFL	0.000.00	0.160.07	0.140.04	0.700.28	0.990.20	2.850.54
	VACA	13.140.73	3.820.01	3.720.05	61.855.06	49.314.11	92.067.93
5-CHAIN LIN	Causal NF	0.010.00	0.120.02	0.080.01	0.620.19	0.690.15	1.910.44
	CAREFL	0.000.00	0.470.23	0.460.22	0.790.41	1.190.25	4.210.87
	VACA	17.310.84	5.950.05	6.060.08	103.7510.04	80.8111.06	124.5220.86
COLLIDER LIN	Causal NF	0.000.00	0.020.01	0.010.00	0.460.12	0.560.11	0.950.19
	CAREFL	0.000.00	0.020.01	0.010.00	0.390.07	0.450.05	0.740.07
	VACA	13.450.43	0.220.01	0.860.02	37.223.55	28.774.22	71.216.73

Fairness Outcomes: Accuracy and Counterfactual Unfairness

- We train classifiers on exogenous u (with sensitive parts removed) and achieve zero counterfactual unfairness.
- We match or exceed the accuracy/F1 of unaware and fair- x baselines.
- We demonstrate practical value of causal NFs for fair decision-making.

Table: Accuracy, F1-score, and counterfactual unfairness of audited classifiers.

	Logistic classifier				SVM classifier			
	full	unaware	fair x	fair u	full	unaware	fair x	fair u
f1	72.286.16	72.374.90	59.668.57	73.084.38	76.042.86	76.805.82	68.285.74	77.391.52
accuracy	67.003.83	66.752.63	54.755.91	66.503.70	69.503.11	71.003.83	59.252.99	69.751.26
unfairness	5.842.93	2.810.72	0.000.00	0.000.00	6.652.45	2.780.40	0.000.00	0.000.00

Fairness Use Case: Partial Graph with SCC Blocks

- We model with partial causal knowledge using SCC grouping of discrete variables.
- We learn ANF within SCCs and preserve inter-block edges; we enable fair interventions on sensitive features.
- We build a foundation for counterfactually fair classification using exogenous variables.




Figure: German Credit: SCC blocks (rectangles), inter-SCC causal edges (solid), and ANF learning order (dashed).

Open Questions

- Can we handle hidden confounding with limited interventional data while preserving identifiability?
- How do we relax bijectivity/acyclicity (e.g., cycles, non-bijective mechanisms) without losing causal guarantees?
- Can flow-matching or other training objectives improve data efficiency while maintaining causal consistency?
- What guarantees can we provide under partial graphs and imperfect dequantization in mixed data?
- How do we best deploy causal NFs for fair decision-making and discovery in large domains (e.g., neuroimaging)?

- Thank you for your attention!
- Questions and feedback are welcome.

- We thank collaborators and the broader causal ML community for foundational ideas.
- Open-source code available to support reproducibility and adoption.



Papamakarios, G., Pavlakou, T., and Murray, I. 2017. Masked Autoregressive Flow for Density Estimation. *Advances in Neural Information Processing Systems (NeurIPS)*, 30.



Durkan, C., Bekasov, A., Murray, I., and Papamakarios, G. 2019. Neural Spline Flows. *Advances in Neural Information Processing Systems (NeurIPS)*, 32.