


Challenges in Training PINNs: A Loss Landscape Perspective

Pratik Rathore; Weimu Lei; Zachary Frangella; Lu Lu; Madeleine Udell

01 Title and Attribution


- Pratik Rathore 1; Weimu Lei 2; Zachary Frangella 3; Lu Lu 4; Madeleine Udell 2 3
- Focus on **loss landscape** challenges and [practical improvements](#).

02 Abstract: Tackling Ill-Conditioning in PINN Optimization



- PINN training failures stem from operator-induced ill-conditioning; Adam plus L-BFGS helps; NysNewton-CG improves further.
- Theory links **operator conditioning** to **loss conditioning**.
- Hybrid first and second order methods deliver [lower loss](#) and [lower error](#).


03 Introduction: Ill-Conditioning in PINNs and Hybrid Second-Order Solutions



- Differential operators make PINN losses ill-conditioned; quasi and second order phases fix curvature.
- Hessian spectra visualizations reveal **severe spread** and near zero mass.
- A hybrid schedule achieves [faster convergence](#) and [better accuracy](#).

04


Surprising result: L-BFGS preconditioning cuts condition number $\sim 1000\times$, unlocking near-zero loss.



- A short L-BFGS phase collapses condition numbers by about one thousand times, unlocking near zero residuals.
- Preconditioning exposes and fixes **hard curvature** hidden from first order progress.
- Spectra compress and residual maps become **uniformly dim**.
- Holds across convection, reaction, and wave PDEs; widths 50–400.

05


Problem Setup: PINN Formulation and Experimental Protocol



- Nonlinear least squares over residual and boundary/initial terms with multiple PDEs and widths using Physics-Informed Neural Networks (PINNs).
- Architectures are 3-layer tanh networks; metrics include L2 relative error (L2RE).
- Optimizers compared include **Adam**, **L-BFGS** (strong Wolfe conditions), and the hybrid sequence.

06

Related Work: PINNs vs Operator Learners and Tackling Ill-Conditioning



- Physics-Informed Neural Networks (PINNs) enforce physics without solution data, while operator learners may violate physics.
- Prior tactics adjust losses or architectures but leave **right preconditioning** underexplored.
- Our analysis complements kernel-based views with **finite-width spectra**.

07


Title

Qualitative Setup and Walkthrough

- Convection, reaction, and wave problems illustrate distinct artifact patterns.
- Training failure shows streaking, ringing, and residual hotspots near fronts and boundaries.
- Successful runs yield **smooth residuals** and **flat error fields**.
- Convection: bands $\rightarrow$ pockets $\rightarrow$ smooth (Adam $\rightarrow$ L-BFGS $\rightarrow$ NNCG).
- Reaction: lag/ringing $\rightarrow$ damped $\rightarrow$ clean transient.
- Wave: boundary ringing/checkerboard $\rightarrow$ corrected phase $\rightarrow$ aligned amplitude.
- Convection: front-aligned bands $\rightarrow$ pocketed specks $\rightarrow$ featureless field.
- Reaction: overshoot/lag $\rightarrow$ damped ripples $\rightarrow$ smooth stiff region passage.
- Wave: edge ringing/checkerboard $\rightarrow$ phase-corrected $\rightarrow$

Qualitative Walkthrough by PDE: Convection, Reaction, Wave

- Convection: Adam leaves banded fronts and residual hotspots; L-BFGS sharpens and reduces bands; NNCG removes fine ripples.
- Reaction: Adam shows overshoot/undershoot and phase lag; L-BFGS damps ringing and aligns the stiff transient; NNCG smooths the passage.
- Wave: Adam has boundary ringing and checkerboard speckle; L-BFGS removes checkerboard and improves phase; NNCG aligns amplitude and phase.
- Convection — phase cues: Adam $\Rightarrow$ banded, front-aligned errors; L-BFGS $\Rightarrow$ small pockets; NNCG $\Rightarrow$ smooth, uniformly dim fields.
- Reaction — phase cues: Adam $\Rightarrow$ lag/oscillations; L-BFGS $\Rightarrow$ damped ringing and closer stiff segment tracking; NNCG $\Rightarrow$ clean transient.
- Wave — phase cues: Adam $\Rightarrow$ boundary ringing/checkerboard; L-BFGS $\Rightarrow$ phase correction; NNCG




- Lower residual loss strongly correlates with lower L2 relative error across problems.
- High accuracy optimization is **necessary** for good predictions.
- Occasional mismatches motivate checking **error fields** directly.

- At loss $\sim 1e^{-123}$: artifacts align with fronts/boundaries; bands and ringing are visible.
- At loss $\sim 1e^{-125}$: artifacts vanish; error fields flatten by $\sim 10^7$.
- Residual loss vs L2RE trend is tight and monotone; large gains appear in the low-loss tail.
- Convection @ $1e^{-123}$: wavy, front-aligned errors and boundary ringing remain.
- Convection @ $1e^{-125}$: sharp front, featureless interior; errors uniformly dim.

Convection: fronts go from wavy/banded to sharp and smooth; residual hotspots shrink to specks.


Reaction: overshoot/lag with ringing becomes a clean pass through the stiff transient.

Wave: boundary ringing/checkerboard fades; amplitude and phase align across the field.




- Blue histogram = Hessian eigenvalue density; orange = after L-BFGS preconditioning.
- Left panels: residual vs boundary/initial contributions.
- Compressed spread $\Rightarrow$ easier curvature; near-zero mass $\Rightarrow$ flat directions.
- Spectra show large outliers and mass near zero indicating severe ill conditioning.
- Maximum eigenvalues reach thousands to hundreds of thousands depending on the problem.
- These features predict **slow first order** progress near optimality.

12 Residual Drives Ill-conditioning: Component-wise Intuition




- Residual terms dominate conditioning compared to boundary and initial components.
- Differential operator properties transfer to **loss curvature**.
- Preconditioning targets the residual induced **bottleneck**.

13 L-BFGS as Preconditioner: How It Changes the Spectrum



- Eigenvalue magnitudes and condition numbers compress by about one thousand times.
- Local curvature improves and search directions become **more effective**.
- Preconditioning benefits appear across problems and loss components.

14 Adam+L-BFGS: Empirical Superiority and Why It Works




- Adam+L-BFGS reaches the smallest minimum and median losses across widths.
- Adam explores saddles; L-BFGS conditions curvature for **clean residuals**.
- Single methods leave persistent structures or plateaus **unresolved**.
- Why it works: explore with Adam, then condition and clean up with L-BFGS.

15 Qualitative Switch Points: What the Solution Looks Like at Each Phase

- After Adam: BC/IC (boundary/initial conditions) largely satisfied; interior shows coherent residual structures (front-aligned bands, ringing).
 - Vivid cue: banded fronts and ringing persist; structured artifacts dominate.
 - Extra cue: front-aligned bands dominate; boundary ringing is pronounced.
 - Concise cue: front-aligned bands prevail across the field.
- After L-BFGS: bands collapse into localized pockets; curvature is preconditioned; residual amplitude drops broadly.
 - Vivid cue: bands $\rightarrow$ small pockets; phase alignment visibly improves.
 - Extra cue: sharper fronts; pockets concentrate at a few hard regions.
 - Concise cue: bands collapsed into small pockets.
- After NNCG: fine ripples damp out; gradients fall further; fields look uniformly dim with smoother transitions.
 - Vivid cue: uniformly dim fields with smooth transitions; pockets dissolve.
 - Extra cue: featureless field with flat residual/error maps.
 - Concise cue: featureless, uniformly dim field.

16 NysNewton-CG (NNCG): Idea and When to Use




- After Adam plus L-BFGS, use damped Newton conjugate gradients (CG) with Nyström (approximation) preconditioned CG to reduce loss and gradients.
- It addresses Wolfe conditions line-search stalls that cause **zero steps**.
- It improves L2 relative error (L2RE) beyond the hybrid baseline when fine ripples persist.

17 NNCG Outcomes: Qualitative Improvements Over Adam+L-BFGS


- Error and residual maps clean further with noticeable **gradient drops**.
- L2 relative error improves across convection, reaction, and wave.
- Gains appear after line-search stalls in the quasi-Newton phase.

18 Practical Recipe: Adam $\rightarrow$ L-BFGS $\rightarrow$ NNCG Schedule



- Trigger phase switches from residual and error visuals and line search behavior rather than iteration counts.
- Warm start with Adam, switch to L-BFGS on persistent bands, and fine tune with NNCG when faint ripples remain.
- Stop when residual maps are featureless and gradients are **consistently small**.


19 From Operator Ill-conditioning to PINN Loss Conditioning



- Ill conditioned operators imply ill conditioned PINN losses under reasonable training assumptions.
- First order methods require many steps proportional to an effective condition number.
- Curvature aware steps enable **faster high accuracy** progress.

20

Theory Takeaways I: Why First-order Slows Near Optimum



- Condition numbers increase with sample size causing **slower first order** convergence.
- Damped Newton achieves linear convergence after a suitable warm start.
- The analysis supports the Adam to L-BFGS to NNCG schedule in practice.

21

Algorithm Intuition: GDND Phases and Convergence Picture

- Phase one uses gradient based warm start to reach a good neighborhood.
- Phase two applies damped Newton with line search to obtain linear convergence.
- L-BFGS approximates Newton efficiently before NNCG refines the solution **more precisely**.

22

Limitations and When This Might Fail

Optimizer	Convection	Reaction	Wave
L-BFGS	4.6e-2	3.6e-2	9.0e-2
NNCG	2.5e-1	7.2e-1	2.9e1
Time Ratio	5.43	20	322.22

- Residual only training can mis rank solutions, so inspect error fields directly.
- NNCG is costlier per iteration and should be reserved for fine tuning.
- Extreme stiffness or sparse sampling may require adaptive strategies to avoid **failures**.

23

Conclusion: Conditioning Drives PINN Training; Second-Order Methods Resolve It

Optimizer	Convection		Reaction		Wave	
	Loss	L2RE	Loss	L2RE	Loss	L2RE
Adam	1.40e-4	5.96e-2	4.73e-6	2.12e-2	2.03e-2	3.49e-1
L-BFGS	1.51e-5	8.26e-3	8.93e-6	3.83e-2	1.84e-2	3.35e-1
Adam+L-BFGS	5.95e-6	4.19e-3	3.26e-6	1.92e-2	1.12e-3	5.52e-2

- Hybrid first and second order training with NNCG fine tuning yields the best accuracy.
- Operator induced ill conditioning explains training stalls and optimization **failures**.
- Preconditioning and Newton CG address curvature and deliver **robust improvements**.

Thank You

Questions & Discussion