

Foundation Model Insights and a Multi-Model Approach for Superior Fine-Grained One-shot Subset Selection

Zhiqing Wan, Zhixiang Wang, Zheng Wang, Xin Xu, Shin'ichi Satoh

National Engineering Research Center for Multimedia Software, Institute of Artificial Intelligence, School
of Computer Science, Wuhan University, China
Hubei Key Laboratory of Multimedia and Network Communication Engineering
The University of Tokyo, Japan
National Institute of Informatics, Japan
School of Computer Science and Technology, Wuhan University of Science and Technology, Wuhan
430065, China

May 24, 2026

Zhijing Wan et al. (Wuhan Univ.)

Foundation Model Insights..

May 24, 2026

1 / 25

Why We Study One-Shot Subset Selection

- We see that training on full datasets is costly in compute, time, and energy.
- We can preserve accuracy with smaller, well-chosen subsets while cutting cost.
- We leverage Foundation Models (FMs) for strong, general-purpose representations without task-specific pretraining.

Zhijing Wan et al. [Wuhan Univ.]

Foundation Model Insights

May 24, 2026

2 / 2

Our Problem: One-Shot Subset Selection with Minimal Compute

- Our goal: we pick a small, informative subset once (no iterative retraining).
- We note that traditional pipelines rely on dataset-pretrained IEs, which are costly and dataset-dependent.
- We ask: Do FMs improve selection, and are all FMs equally good as IEs?

Zhijing Wan et al. (Wuhan Univ.)

Foundation Model Insights...

May 24, 2026

3 / 20

From Dataset-Dependent IEs to Foundation Models

- We replace task-specific IEs with FMs to remove dataset dependence.
- We investigate when FMs help and which FM we should use.




Figure: One-shot subset selection: traditional IE vs. a single FM. Our FM-based IE excels on fine-grained data.

Zhijing Wan et al. [Wuhan Univ.]

Foundation Model Insights.

May 24, 2026

4/2

- We find that FMs shine on fine-grained data; their advantage fades on noisy, coarse-grained data.
- We show that higher FM task accuracy does not imply better subset selection.
- We propose RAM-APL, which fuses intra-class prototypicality (RAM) and inter-class ambiguity (APL) with adaptive weighting.
- We achieve state-of-the-art results on Pet, Food-101, and CUB-200-2011.

RAM: Intra-class Prototypicality via Ranking

- We compute per-FM class centers and rank samples by distance (closer means more prototypical).
- We normalize ranks to align unaligned feature spaces across FMs.
- We average ranks across FMs to obtain a robust prototypicality score; lower scores indicate more representative samples.

- We extract features per sample from multiple FMs (no target-dataset finetuning).
- We compute RAM: we rank within each class by distance to the class center per FM; we then average normalized ranks across FMs.
- We compute APL: we use pseudo-label correctness across FMs to capture inter-class ambiguity.
- We fuse scores with sampling-rate-dependent weights and we select the smallest fused scores.

RAM in Action: What Low vs. High RAM Looks Like

- We observe that low-RAM samples have clear objects, simple backgrounds, and prototypical views.
- We observe that high-RAM samples exhibit atypical poses, occlusions, or cluttered backgrounds.



Figure: Qualitative examples illustrating the RAM metric.

- We assign a pseudo-class by nearest class center for each FM.
- We mark correctness against ground truth and average across FMs.
- We treat lower APL (more confusion) as flags for discriminative, boundary samples.

- $\text{Score} = W1 \times \text{RAM} + W2 \times (1 - \text{APL})$.
- At low budget p , we prioritize $W1$ (representative/easy).
- As p grows, we increase $W2$ to include more discriminative/hard samples.
- By default, we use alpha 0.2 and beta 1 to provide robust weighting across datasets.

Why Multi-FM? Complementarity of Representations

- We find that different FMs encode diverse semantics; naïve “best FM” heuristics fail.
- We observe low cosine similarity across FMs, which justifies fusion over a single FM.
- We use ranking-space fusion to avoid alignment and concatenation pitfalls.




Figure: Cosine similarity matrix across FM embeddings.

Experimental Setup: Datasets, Baselines, Protocol

- We evaluate on CIFAR-10/100, CIFAR-10N/10I, Oxford-IIIT Pet (and a noisy variant), CUB-200-2011, and Food-101.
- We compare to baselines: diversity—K-Center Greedy/CoreSet [Sener18], MDS (Kennard–Stone) [Kennard69], CD (Cluster-Distance, k-means) [MacQueen67]; graph-cut—GC [Wei13]; uncertainty—Margin [Settles10]; optimization—GraNd [Paul21], GLISTER [Killamsetty21], CRAIG [Mirzasoleiman20], CAL (class-balanced CRAIG) [Mirzasoleiman20]; herding—Herding [Chen12]; prototype—MIN (Nearest-Prototype; prototypical proximity) [Snell17]; plus Random.
- We use MobileNet-V3 as one of the target models and evaluate across multiple sampling rates.
- We conduct a Single-Model Study and multi-model (RAM-APL) experiments.

Single-Model Study Protocol

- We compare three IE types: model-TD (target-pretrained), model-TIN, and FMs.
- We vary selection algorithms and sampling rates.
- We diagnose when FMs help and which FM is best for selection.




Figure: Framework of our Single-Model Study.

Finding 1: Fine-Grained Gains; Coarse/Noisy Are Mixed

- On fine-grained datasets (e.g., Pet, CUB-200-2011), FMs are frequently optimal across budgets and selectors.
- On clean but coarse-grained data (e.g., CIFAR-100), FM advantage is less consistent.
- Under label noise, FMs remain strong for fine-grained Pet but offer limited gains on noisy coarse-grained CIFAR-10.




Figure: FM-based selection excels on fine-grained tasks; benefits are mixed for coarse or noisy data.

Finding 2: FM Task Accuracy Does Not Imply Selection Quality

- We find that higher FM zero-shot or finetuned accuracy does not ensure better selection.
- We observe that different FMs win under different budgets and selection algorithms.
- We are motivated to use multi-FM fusion rather than pick a single “best” FM.




Figure: FM target accuracy does not equal subset selection quality across models.

We Outperform Baselines on Fine-Grained Datasets

- We achieve consistent gains over Random, GC [Wei13], MDS [Kennard69], and KCG [Sener18] across budgets.
- We evaluate on Oxford-IIIT Pet, Food-101, and CUB-200-2011.
- We find that adaptive fusion selects representative first and then discriminative samples effectively.




Figure: Comparison with baselines on three classic fine-grained datasets (average accuracy).

Baselines: Herding [Chen12]; K-Center Greedy/CoreSet [Sener18]; GC (Graph Cut) [Wei13]; MDS (Kennard-Stone) [Kennard69]; CD (k-means cluster distance) [MacQueen67]; Margin [Settles10]; Forgetting [Toneva19]; GraNd [Paul21]; GLISTER [Kilamsetty21]; CRAIG

Ablation: RAM vs. RAM-APL vs. MIN (Pet) — Mandatory Table 1

- We find that RAM (multi-FM ranks) already surpasses single-FM MIN [Snell17].
- We observe that adding APL further improves performance, especially at higher budgets.
- We conclude that combining intra-class and inter-class signals helps.

Table: Table 1. Ablation study on Pet. “Model-TD” refers to the model pretrained on Pet for 10 epochs. Values are Top-1 accuracy (mean±std).

Method	IE	1%	50%	70%
Model-TD	Model-TD	5.6±0.7	40.3±2.6	55.2±2.7
MIN [Snell17]	CLIP	5.6±0.2	45.9±1.8	56.3±0.7
MIN [Snell17]	DINOv2	6.2±0.1	46.8±2.0	60.5±2.9
RAM (ours)	CLIP+DINOv2	5.9±0.3	47.1±1.4	56.5±2.7
RAM-APL (ours)	CLIP+DINOv2	6.5±0.4	47.5±1.9	58.7±2.2

Fusion Matters: Ranking-Space Fusion vs. Concatenation

- We find that simple concatenation underperforms our RAM-APL ranking-space fusion.
- We observe that ranking-space fusion is robust to FM misalignment and scale differences.
- We see consistent benefits across sampling rates.

Table: Table 3. Comparison of feature fusion strategies (Top-1 accuracy, mean±std).

Fusion strategy	1%	10%	30%	50%	70%
Concatenate	5.9±0.4	16.3±0.4	31.7±1.3	47.7±3.0	57.8±1.2
Ours (RAM-APL)	6.5±0.4	15.2±1.2	32.4±2.9	47.5±1.9	58.7±2.2

Benchmark Comparison: Food-101

- We improve Top-1 over Random by an average of +4.44%.
- We consistently outperform strong baselines such as GC [Wei13], MIN [Snell17], and Forgetting [Toneva19].
- We scale to large fine-grained datasets with strong gains at low budgets.

Table: Table 9. Comparison with baselines on Food-101. “Traditional” stands for models pretrained on Food-101 for ten epochs. Means and variances are across three runs. Best in bold; second-best in blue. Tags indicate sources.

Methods	IE	1%	10%	30%	50%	70%	Avg. vs. Random
Random	-	6.8±0.4	45.5±0.5	69.5±0.2	76.2±0.3	79.6±0.3	-
Herdg [Chen12]		6.6±0.5	42.6±1.2	64.6±0.2	74.3±0.2	78.7±0.3	-2.16
KCG [Sener18]		4.7±0.3	35.0±0.7	67.2±0.2	75.9±0.2	79.6±0.2	-3.04
CD (cluster diversity) [MacQueen67]		3.9±0.1	27.7±0.6	66.2±0.4	75.7±0.2	80.0±0.1	-4.82
Margin [Settles10]		5.0±0.7	36.8±2.7	68.4±0.3	76.4±0.1	79.8±0.2	-2.24
Forgetting [Toneva19]		9.5±0.5	53.4±0.5	71.7±0.1	77.3±0.1	79.9±0.1	+2.84
GraNd [Paul21]		2.9±0.2	18.4±0.4	59.7±0.5	74.7±0.2	80.0±0.1	-8.38
CAL/CRAIG [Mirzasoileman20]		10.3±0.4	51.8±0.3	69.5±0.5	76.0±0.1	79.5±0.2	+1.90
Gister [Killamsetty21]		6.9±0.3	38.8±1.1	67.4±0.4	75.9±0.4	80.1±0.2	-1.70
GC [Wei13]		10.8±0.4	54.4±0.3	71.3±0.4	76.8±0.2	79.5±0.3	+3.04
MDS [Kennard69]		6.1±0.5	45.1±0.4	69.7±0.3	76.2±0.2	79.6±0.1	-0.18
MIN [Snell17]		9.4±0.2	52.1±0.6	70.7±0.7	76.7±0.1	79.8±0.1	+2.22
Ours (RAM-APL)	CLIP+DINOv2	12.3±0.3	56.1±0.6	72.3±0.1	77.9±0.1	81.2±0.2	+4.44

Baselines: Herding [Chen12]; K-Center Greedy/CoreSet [Sener18]; GC (Graph Cut) [Wei13]; MDS (Kennard-Stone) [Kennard69]; CD (k-means cluster distance) [MacQueen67]; Margin [Settles10]; Forgetting [Toneva19]; GraNd [Paul21]; CAL/CRAIG [Mirzasoileman20]; Gister [Killamsetty21].

Generalization Across Target Architectures

- We show that selected subsets transfer to different target models (e.g., MBV3).
- We observe that RAM-APL surpasses baselines even when the IE differs from the target.
- We support practical, model-agnostic subset selection.

Table: Table 6. Cross-architecture generalization of RAM-APL. Target model: MobileNet-V3 (MBV3).

Method	IE → Target Model	10%	30%	50%
Random	MBV3 → MBV3	10.9±1.1	42.1±3.6	61.6±1.9
Forgetting [Toneva19]	MBV3 → MBV3	13.3±0.8	42.0±2.0	61.0±2.3
GC [Wei13]	MBV3 → MBV3	12.4±1.7	40.4±0.2	61.3±1.5
MDS [Kennard69]	MBV3 → MBV3	11.9±0.7	39.8±2.0	62.1±3.3
MIN [Snell17]	MBV3 → MBV3	11.9±1.8	38.4±0.8	61.0±1.4
RAM-APL (Ours)	(CLIP+DINOv2) → MBV3	13.6±0.3	45.7±0.9	62.3±1.4

Baselines: Herding [Chen12]; K-Center Greedy/CoreSet [Sener18]; GC (Graph Cut) [Wei13]; MDS (Kennard-Stone) [Kennard69]; CD (k-means cluster distance) [MacQueen67]; Margin [Settles10]; Forgetting [Toneva19]; GraNd [Paul21]; CAL/CRAIG [Mirzasoileman20]; Gister [Killamsetty21].

- We find that setting alpha to 0.2 and beta to 1 outperforms an equal-weight fusion.
- We support a curriculum-style weighting that increases the emphasis on ambiguity as the budget grows.
- We observe robust default behavior across datasets.




Figure: Parameter analysis at 70% sampling on Pet: our scheduled fusion outperforms an equal-weight, no-schedule fusion of RAM and the complement of APL.

Future Work & Broader Impact

- We plan automatic FM weighting/selection and better fusion for coarse or noisy settings.
- We aim to extend beyond images (e.g., audio, video, multimodal) and beyond classification.
- We will integrate robustness and fairness constraints to protect minorities and rare cases.
- We enable scalable training on large or evolving datasets with lower cost.

- We see limited advantage on noisy, coarse-grained datasets (e.g., CIFAR-10N).
- We note that FM choice depends on dataset, budget, and selector—there is no single winner.
- We may risk overlooking rare-but-important cases without explicit fairness or coverage constraints.

References (Baselines)

- Chen, Y., Welling, M., Smola, A., & Li, K. (2012). Super-samples from Kernel Herding. UAI. [Chen12]
- Sener, O., & Savarese, S. (2018). Active Learning for Convolutional Neural Networks: A Core-Set Approach. ICML. [Sener18]
- Kennard, R. W., & Stone, L. A. (1969). Computer Aided Design of Experiments. Technometrics. [Kennard69]
- Wei, K., Iyer, R., & Bilmes, J. (2013). Fast Multi-Document Summarization via Submodular Maximization. ACL. [Wei13]
- Settles, B. (2010). Active Learning Literature Survey. Tech. Report, UW–Madison. [Settles10]
- Toneva, M., Sordoni, A., Niculae, V., et al. (2019). An Empirical Study of Example Forgetting During Deep Neural Network Learning. NeurIPS. [Toneva19]
- Paul, M., Ganguli, S., & Dziugaite, G. K. (2021). Deep Learning on a Data Diet: Finding Important Examples Early in Training. NeurIPS. [Paul21]
- Killamsetty, K., Durga, S., & Iyer, R. (2021). GLISTER: Generalization based Data Subset Selection for Efficient and Robust Learning. AAAI. [Killamsetty21]
- MacQueen, J. (1967). Some methods for classification and analysis of multivariate observations. Proc. 5th Berkeley Symposium on Mathematical Statistics and...

- We thank our collaborators, institutions, and funding agencies.
- We appreciate the open-source FM community for enabling this research.

- How can we automatically select and weight FMs for a given dataset and budget?
- Can we extend RAM-APL beyond images (e.g., audio, video, multimodal) and beyond classification?
- How do we make selection robust to extreme label noise, class imbalance, and domain shift?
- Can we enforce fairness or coverage constraints to protect rare subpopulations without hurting accuracy?
- What are the best low-cost proxies for FM complementarity when scaling to many FMs?