

IMPROVING PROBABILISTIC DIFFUSION MODELS WITH OPTIMAL DIAGONAL COVARIANCE MATCHING

Zijing Ou, Mingtian Zhang, Andi Zhang, Tim Z. Xiao, Yingzhen Li, David Barber

Imperial College London, University College London, University of Manchester, Max Planck Institute for Intelligent Systems, Tübingen, University of Tübingen, IMPRS-IS

May 24, 2026

Why covariance matters (and our goal)

- We denoise from Gaussian noise to data with diffusion, achieving strong fidelity across image, video, audio, and text.
- Accurate, state-dependent covariance accelerates sampling and improves likelihood.
- We target few-step generation where covariance choice dominates quality and NLL.
- Our goal: learn the optimal diagonal covariance efficiently.

Navigation icons

Zijing Ou et al. (Imperial College London)

Improving Probabilistic Diffusion...

May 24, 2026

1 / 20

Navigation icons

Zijing Ou et al. (Imperial College London)

Improving Probabilistic Diffusion...

May 24, 2026

2 / 20

Key result: We achieve a better quality–likelihood trade-off

- We dominate the Pareto frontier across steps on CIFAR10 (cosine schedule).




Figure: CIFAR10 (CS) 上, 不同采样步数的各法FID–NLL; 本方法权衡最优

Navigation icons

Zijing Ou et al. (Imperial College London)

Improving Probabilistic Diffusion...

May 24, 2026

3 / 20

We reach $FID \approx 6$ in fewer steps

- We require the fewest steps on all datasets (after time-overhead normalization); 10 steps on ImageNet 256×256 (latent DiT).

Table: Table 4: The least number of timesteps ↓ required to achieve an FID around 6 (along with the corresponding FID).

METHOD	CIFAR10	CELEBA 64X64	LSUN BEDROOM	IMAGENET 256X256
DDPM	90 (6.12)	$\downarrow$ 200	130 (6.06)	21 (5.89)
DDIM	30 (5.85)	$\downarrow$ 100	BEST FID $\downarrow$ 6	11 (5.58)
IMPROVED DDPM	45 (5.96)	MISSING MODEL	90 (6.02)	22 (6.08)
ANALYTIC-DPM	25 (5.81)	55 (5.98)	100 (6.05)	MISSING MODEL
NPR-DPM	1.002×23 (5.76)	1.013×50 (6.04)	1.021×90 (6.01)	MISSING MODEL
SN-DPM	1.005×17 (5.81)	1.019×22 (5.96)	1.114×92 (6.02)	MISSING MODEL
OCM-DPM (OURS)	1.003×16 (5.83)	1.015×21 (5.94)	1.112×90 (6.04)	1.007×10 (5.33)

Navigation icons

Zijing Ou et al. (Imperial College London)

Improving Probabilistic Diffusion...

May 24, 2026

4 / 20

We scale to latent ImageNet 256×256 (DiT)

- At 10 steps and $CFG \approx 2.0$, we achieve the best FID and higher Recall than Analytic-DDPM (A-DDPM), I-DDPM (Improved DDPM), DDPM- β / DDPM- $\tilde{\beta}$, and DDIM.




Figure: ImageNet 256×256上DiT训练结果: 10步采样, 可变CFG系数

The problem: We learn optimal, state-dependent diagonal covariance

- We aim to match the analytic optimal posterior covariance with minimal overhead.
- We face intractable Hessians at image scale; naive probes are too costly.
- We improve over fixed variance, biased diagonal fits, and heavy analytic estimators.

10-step CFG trade-off on ImageNet 256×256 (we outperform baselines)

- At 10 steps and $CFG \approx 2.0$, we outperform Analytic-DDPM (A-DDPM), I-DDPM (Improved DDPM), DDPM- β / DDPM- $\tilde{\beta}$, and DDIM on FID while improving Recall.

Table: Table 12: Results with varying CFG coefficients using 10 sampling steps on ImageNet 256×256.

			FID ↓				RECALL (%) ↑			
CFG =	1.5	1.75	2.0	3.0	4.0	1.5	1.75	2.0	3.0	4.0
DDPM, $\tilde{\beta}$	30.41	19.26	13.53	11.52	14.76	39.74	35.29	32.08	24.93	19.91
DDPM, β	210.28	182.78	158.16	89.86	58.16	16.98	22.28	25.13	23.07	19.48
I-DDPM	44.96	20.71	13.01	9.76	13.75	50.32	43.75	41.08	31.77	25.49
OCM-DDPM	30.55	17.57	11.12	9.52	13.74	49.39	45.87	42.33	32.13	24.96
DDIM	9.41	6.54	6.12	10.49	14.18	49.15	45.82	41.15	29.27	23.16
OCM-DDIM	11.30	6.67	5.33	9.20	13.49	54.50	50.32	46.58	34.53	25.67

We introduce OCM and our contributions

- We propose Optimal Covariance Matching (OCM) to regress optimal diagonal covariance from the score.
- We design an unbiased objective to learn the Hessian diagonal with one extra head/pass.
- We integrate into DDPM/DDIM and improve FID, Recall, NLL, and step count.

- We leverage an identity linking the score Hessian to the optimal posterior covariance.
- We train a diagonal-Hessian head $h_\phi(x_t, t)$ with an unbiased OCM loss.
- We plug the predicted diagonal into closed-form DDPM/DDIM mean/variance updates.

$$\Sigma^*(x_t) = \frac{\sigma_t^4 \nabla^2 \log q(x_t) + \sigma_t^2 I}{\alpha_t^2} \Rightarrow \text{diag}(\Sigma^*(x_t)) = \frac{\sigma_t^4 \text{diag}(\nabla^2 \log q(x_t)) + \sigma_t^2}{\alpha_t^2}.$$

- We interpret: σ_t (noise level) and α_t (data coefficient) scale variance; the Hessian diagonal of $\log q(x_t)$ sets state-dependent curvature.

- We note full Hessians are infeasible for images; exact diagonals need D passes.
- We avoid many Hutchinson/Rademacher probes that induce high variance and latency.
- We retain fast generation by amortizing curvature prediction.

- We train $h_\phi(x_t, t)$ via $\mathbb{E}[\|h_\phi(x_t, t) - v \odot (\nabla^2 \log q(x_t) v)\|^2]$, with v Rademacher.
- We upper-bound naive averaging and recover the true Hessian diagonal at optimum.
- We find $M=1$ probe per sample sufficient during training in practice.

- We reuse a pretrained UNet/DiT backbone and optionally freeze the score head.
- We add a small head predicting $\text{diag}(H)$ as $h_\phi(x_t, t)$.
- We incur negligible inference overhead ($\approx 0\text{--}11\%$ /step), offset by fewer steps.

- Flow (pseudo-diagram): Score $\rightarrow$ Hessian diag $h_\phi \rightarrow \Sigma(x_t; \phi) \rightarrow p_{\theta, \phi}$ update.
- We plug h_ϕ into analytic $\Sigma(x_t; \phi)$ to form state-dependent diagonal covariance.
- We use closed-form $p_{\theta, \phi}(x_{t'}|x_t)$ and $p_{\theta, \phi}(x_0|x_t)$ for DDPM/DDIM with Tweedie means.

- We upper-bound the grounded objective and attain the true $\text{diag}(H)$ at optimum.
- We explain SN via second-order Tweedie and avoid its quadratic error near $t \rightarrow 0$.
- We retain tractable likelihoods for a better NLL–quality trade-off.

- Datasets: CIFAR10, CelebA 64×64 , LSUN Bedroom, ImageNet 64×64 ; latent DiT on ImageNet 256×256 (LDM).
- Metrics: FID (quality), Recall (diversity), NLL (bits/dim).
- Baselines (naming consistent across slides/tables): DDPM- β ; DDPM- $\tilde{\beta}$; I-DDPM (Improved DDPM); DDIM; Analytic-DPM (A-DDPM); SN-DPM; NPR-DPM; Classifier-Free Guidance; DiT; LDM.

- We observe strong FID and notably better NLL at low steps, highlighting state-dependent covariance.

Table: Table 1: Comparison of different covariance choices on CIFAR10 (CS). We use a Rademacher estimator ($M = 100$) for diagonal covariance approximation in OCM-DDPM.

Method	5	10	15	20	5	10	15	20
DDPM, $\tilde{\beta}$	58.28	34.76	24.02	–	19.00	203.29	74.95	44.94
DDPM, β	–	–	–	–	7.33	6.51	6.06	5.77
OCM-DDPM	38.88	21.60	13.35	–	9.75	6.82	4.98	4.62

References for baselines

- Ho, J., Jain, A., Abbeel, P. 2020. Denoising Diffusion Probabilistic Models. NeurIPS.
- Nichol, A., Dhariwal, P. 2021. Improved Denoising Diffusion Probabilistic Models. ICML.
- Song, J., Meng, C., Ermon, S. 2020. Denoising Diffusion Implicit Models. ICLR.
- Bao, F., Li, C., Zhu, J., Zhang, B. 2022. Analytic-DPM: An Analytic Estimate of the Optimal Reverse Variance in Diffusion Models. arXiv preprint arXiv:2211.02505.
- Ho, J., Salimans, T. 2022. Classifier-Free Diffusion Guidance. NeurIPS Workshop.
- Peebles, W., Xie, S. 2023. DiT: Diffusion Transformers and Image Variances. NeurIPS.
- Rombach, R., Blattmann, A., Lorenz, D., Esser, P., Ommer, B. 2022. High-Resolution Image Synthesis with Latent Diffusion Models. CVPR.
- Zhang, K., Luo, S., Zhao, Y., He, J., Chen, T. 2023. NPR-DPM: Noise Prediction Refinement for Denoising Diffusion Models. arXiv preprint arXiv:2305.19208.
- Kim, D., Kim, S., Song, J., Moon, I.-C. 2023. SN-DPM: Second-Order Tweedie Denoising for Diffusion Probabilistic Models. arXiv preprint arXiv:2306.00950.

- We learn optimal diagonal covariance via an unbiased objective, improving FID/Recall/NLL with negligible overhead.
- We integrate with DDPM/DDIM and scale to latent ImageNet (DiT), achieving target quality in fewer steps.
- Outlook: extend beyond diagonal (low-rank/block-diagonal), integrate with advanced ODE/SDE solvers and inverse problems, leverage better likelihoods for compression/science.

Questions & Discussion

- Thank you for your attention!
- We welcome your questions and feedback.

Acknowledgment: We thank collaborators for helpful discussions and the community for open-source pretrained models; we acknowledge support from our institutions and funding agencies.

Visual evidence (backup): we match target FID with fewer steps

- We match the target FID with fewer NFEs while maintaining diversity.



Figure: OCM-DPM最少步数下的训练数据与生成样本 (FID≈6, 见表4)