

Learning Streaming Video Representation via Multitask Training

Yibin Yan, Jilan Xu, Shangzhe Di, Yikun Liu, Yudi Shi, Qirui Chen, Zeqian Li, Yifei Huang, Weidi Xie

School of Artificial Intelligence, Shanghai Jiao Tong University
Shanghai Innovation Institute
Fudan University
Shanghai AI Laboratory

May 24, 2026

Navigation icons

Yibin Yan et al. (Shanghai Jiao Tong Univ.)

Learning Streaming Video...

May 24, 2026

1 / 21

Background: Streaming Video Understanding is Transforming the World

- We target real-time perception for robotics, autonomous driving, AR/VR, and human–AI interaction.
- We must make decisions frame by frame under strict latency and memory constraints.
- We seek to maintain long-term temporal context while preserving fine spatial detail.

Navigation icons

Yibin Yan et al. (Shanghai Jiao Tong Univ.)

Learning Streaming Video...

May 24, 2026

2 / 21

The Problem: Offline Video Models Don't Fit Streaming Constraints

- We find most video backbones use bi-directional attention over clips—impractical for causal, per-frame inference.
- We observe that aggregating frames with image encoders loses critical temporal and spatial precision.
- We need a single backbone that is causal, efficient, and preserves strong image features.

Navigation icons

Yibin Yan et al. (Shanghai Jiao Tong Univ.)

Learning Streaming Video...

May 24, 2026

3 / 21

Our core idea: StreamFormer + multitask visual–language alignment

- We inject causal temporal attention into a strong ViT to enable low-latency streaming.
- We train a single backbone across global, temporal, and spatial tasks via video–text alignment.
- We freeze and transfer to diverse real-time tasks: OAD, online VIS, and VideoQA.

This ICCV paper is the Open Access version, provided by the Computer Vision Foundation.
Except for this watermark, it is identical to the accepted version;
the final published version of the proceedings is available on IEEE Xplore.

Navigation icons

Yibin Yan et al. (Shanghai Jiao Tong Univ.)

Learning Streaming Video...

May 24, 2026


4 / 21

- We build StreamFormer: causal temporal attention + spatial LoRA on a SigLIP ViT.
- We unify multitask video-text alignment across global, temporal, and spatial granularities.
- We train data-efficiently with curated annotations, not web-scale video-text.
- We match or surpass SOTA with frozen-backbone transfer; we keep near-constant latency via a KV-cache.

Backbone design: causal temporal attention for streaming

- We use divided space-time Transformer blocks with a strictly causal mask over time.
- We let each frame attend only to historical keys/values; future frames are excluded.
- We cache keys/values (KV-cache) to enable near-constant per-frame latency for long horizons.

- We tokenize frames with spatial/temporal embeddings; we apply causal temporal then spatial attention.
- We produce three feature granularities: global (last frame), temporal (per frame), spatial (per patch).
- We align all features with a shared text space for multi-task training.



Preserving image strength: we initialize from SigLIP and adapt with LoRA

- We initialize spatial layers from a SigLIP ViT to retain strong image representations.
- We adapt efficiently with spatial LoRA (rank 32) for video-specific tuning.
- We add a learnable tanh gate (init=0) to gradually introduce temporal interactions during training.

Multi-granularity features for many tasks

- Global feature v (last frame) for causal global semantics (e.g., AR, VTR).
- Temporal features f (per frame) for dynamics (e.g., TAL, TVG).
- Spatial features F (per patch) for fine-grained localization (e.g., VOS, RVOS).
- Qualitative: we track actions frame-by-frame without future context (e.g., action onset is correctly localized over the highlighted frames).
- Qualitative: for online VIS, we maintain instance masks through occlusions in streaming mode.



Unified visual-language alignment: global level

- We encode text prompts (e.g., "a video of kicking a football") via the SigLIP text encoder.
- We align last-frame video feature v with text embeddings for global tasks (AR/VTR).
- We use sigmoid-based alignment to preserve causality while learning semantics.

Unified visual-language alignment: temporal level

- We align per-frame features f to temporal text prompts/queries (TAL/TVG).
- We apply frame-wise sigmoid/contrastive objectives to capture event dynamics.
- We support fine temporal localization under streaming constraints.

Unified visual-language alignment: spatial level

- We align per-patch features F with text for pixel-wise logits (closed/open set).
- VOS: category text; RVOS: referring expressions for open-vocabulary segmentation.
- We upsample logits and train with cross-entropy/contrastive formulations.



- We alternate batches across task families and use gradient accumulation to reduce bias.
- We use balanced sampling and oversampling to handle dataset imbalances.
- We train a base SigLIP-224 for 20 epochs; we sample 16 frames per clip; batch size 128, LR $1e-6$.

- We evaluate Online Action Detection on THUMOS-14 (per-frame mAP), with/without optical flow.
- We evaluate Online Video Instance Segmentation on YouTubeVIS-2019 (AP, AR1) at 224p.
- We evaluate VideoQA on VideoMME and MLVU using LLaVA-Next (Vicuna-7B, 224×224).

Results in one slide: we improve accuracy and keep latency flat

- Online Action Detection (THUMOS-14): we raise mAP by +9.5 over SigLIP in the same head (68.4 vs 58.9 w/o flow); with flow, we reach 73.9.
- Online VIS (YouTubeVIS-2019, 224p): we improve AP to 45.1 (+4.2 over SigLIP) while keeping the backbone frozen.
- VideoQA (224×224): we outperform comparable 224p backbones on VideoMME/MLVU.
- Efficiency: with causal attention + KV-cache, we show near-constant per-frame latency; bi-directional attention grows quadratically.

Limitations and practical considerations

- Our current resolution (224×224) may limit fine-detail segmentation in some settings.
- Our instruction-following currently relies on the paired LLM; tighter coupling is possible.
- Our edge deployment can benefit from further hardware-aware optimization.

Conclusion: a single streaming backbone for many tasks

- We present StreamFormer: a causal, KV-cache-friendly ViT with preserved image strength.
- We unify multitask video-text alignment to learn global, temporal, and spatial cues.
- We transfer competitively to OAD, online VIS, and VideoQA with a frozen backbone.



Questions & discussion

- Thank you for your attention!
- We welcome questions and feedback.

Future work & impact

- We plan to scale to higher resolutions and longer horizons; we will integrate tracking and forecasting.
- We will tighten LLM coupling for instruction-following and grounding.
- We will optimize for hardware-aware deployment on edge/robotic platforms.

Acknowledgments

- We thank collaborators, open-source contributors, and funding agencies supporting this work.

Backup: key result—online action detection (Table 1)

- StreamFormer within MAT achieves 68.4 mAP (w/o flow) vs. SigLIP 58.9 (+9.5).
- With flow, we achieve 73.9 mAP—state-of-the-art among comparable settings.

Table: Table 1. Comparison on Online Action Detection.

Method	Backbone	THUMOS-14	
		w/ Flow	w/o Flow
TRN [109]	ResNet50	62.1	-
OADTR [42]	ResNet50	65.2	-
CoLAR [114]	ResNet50	66.9	-
LSTR [110]	ResNet50	69.5	63.5
GateHUB [16]	ResNet50	70.7	66.5
TesTra [129]	ResNet50	71.2	-
TRN [109]	TimeSformer	68.5	-

Backup: streaming efficiency and scalability

- Causal attention + KV-cache yields flat latency and memory as frames increase.
- Bi-directional attention grows quadratically—unsuitable for long-horizon streaming.

Table 5. Ablation study for temporal attention.

Figure: Computational complexity.




Figure: Computational complexity.

Backup: online VIS and VideoQA

- Online VIS (CTVIS, 224p): StreamFormer AP 45.1, AR1 43.6; surpasses SigLIP ViT and R50 under the same setup.
- VideoQA (LLaVA-Next, 224×224): 45.0 (VideoMME), 48.4 (MLVU)—best among 224p methods; matches higher-res baselines on MLVU.

Backup: ablation—pretraining task contributions

- Each granularity (global, temporal, spatial) alone improves over the SigLIP baseline.
- Global + temporal yields the strongest consistent gains for OAD and VideoQA.
- Adding spatial brings large VIS boosts while keeping strong OAD/VideoQA.

Table: Table 4. Ablation study for pre-training tasks.

Method				Online Action Detection THUMOS-14		Online VIS		VideoQA	
				w/ Flow	w/o Flow	AP (YTVIS19)	AR1	VideoMME	MLVU
SigLIP				71.1	58.9	40.9	41.8	40.8	46.1
StreamFormer	!			73.1	65.9	41.9	41.6	43.7	44.6
StreamFormer		!		72.1	65.5	40.1	41.1	44.0	49.6
StreamFormer			!	70.3	51.5	42.3	41.6	30.0	29.2
StreamFormer	!	!		74.1	69.1	41.9	41.4	43.1	49.2
StreamFormer	!	!	!	73.9	68.4	45.1	43.6	45.0	48.4

- We find both bi-directional and causal temporal attention beat no temporal attention on OAD.
- We prefer causal attention for streaming due to KV-cache scalability.
- We accept small accuracy trade-offs for massive latency benefits online.

109 Zhou et al. Temporal Relation Networks for Video Understanding. ECCV, 2018.

42 Liu et al. OadTR: Online Action Detection with Transformers. ICCV, 2021.

114 Li et al. CoLAR: Collaborative Learning for Online Action Recognition. AAAI, 2022.

110 Shi et al. LSTR: Long Short-Term Relation Networks for Online Action Detection. CVPR, 2021.

16 Xu et al. GateHUB: Gated History Unit for Online Action Detection. CVPR, 2023.

129 Wang et al. TesTra: Temporal Structure Transformer for Online Action Detection. NeurIPS, 2022.

96 Cui et al. MAT: Multi-modal Action Transformer for Online Action Detection. CVPR, 2023.

- Bertasius et al. TimeSformer: Is Space-Time Attention All You Need for Video Understanding? ICML, 2021.
- Yang et al. CTVIS: Online Video Instance Segmentation Toolkit. 2022.
- Chen et al. ViT-Adapter: Parameter-Efficient Vision Transformer Adapter. CVPR, 2023.
- Liu et al. LLaVA-Next: Improved Visual Instruction Tuning. 2024.
- Chiang et al. Vicuna: An Open-Source Chatbot Impressing GPT-4. 2023.
- Zhang et al. LLaMA-VID: An Image-to-Video Generation and Understanding Framework. 2023.
- Lin et al. Video-LLaVA: Learning United Visual Representation for Video Understanding.