

Model Immunization from a Condition Number Perspective

Amber Yijia Zheng, Cedar Site Bai, Brian Bullins, Raymond A. Yeh

Department of Computer Science, Purdue University

May 24, 2026

Amber Yijia Zheng et al. (Department of Computer Science, Purdue University) Model Immunization... May 24, 2026 1 / 22

Background: Foundation models and fine-tuning are transforming AI

- Open-weight encoders and vision-language models enable rapid adaptation to new tasks.
- Low-cost fine-tuning (linear probing, few-shot) drives wide deployment across domains.
- Optimization geometry (loss curvature) dictates adaptation speed and stability.
- Safety stakes: easy repurposing can enable harmful use with minimal effort.

Amber Yijia Zheng et al. (Department of Computer Science, Purdue University) Model Immunization... May 24, 2026 2 / 22

Background: Why condition numbers matter for learning speed

- Gradient methods converge faster when the Hessian condition number κ is small.
- Ill-conditioned objectives slow, destabilize, or stall fine-tuning.
- Model features shape Hessians of downstream tasks; thus features control κ .
- Target: steer features so benign tasks are well-conditioned while harmful ones are ill-conditioned.

Amber Yijia Zheng et al. (Department of Computer Science, Purdue University) Model Immunization... May 24, 2026 3 / 22

Problem: Selective model immunization

- Goal: make pre-trained models hard to fine-tune on harmful tasks, yet useful for benign tasks.
- Existing works show empirical effects but lack principled definitions and guarantees.
- Key challenge: selectively degrade optimization geometry only for harmful tasks.
- Need: theory clarifying when immunization is possible and how to enforce it.

Amber Yijia Zheng et al. (Department of Computer Science, Purdue University) Model Immunization... May 24, 2026 4 / 22

- Benign and harmful tasks often share features; naive penalties hurt both.
- Optimization-based defenses must generalize to unknown future probes.
- No agreed metric to quantify “immunized” beyond final accuracy.
- Theoretical feasibility depends on data geometry across tasks.

- They define and analyze immunization via Hessians for linear probing through a condition-number framework.
- They introduce a κ -maximizing regularizer to ill-condition harmful tasks.
- They present an end-to-end algorithm balancing κ increase (harmful) and κ decrease (pre-training).
- They report strong validation on linear models and deep nets with large RIR gains.

- Setting: linear feature extractor θ ; squared-loss linear probes on tasks.
- Task Hessian takes form $H(\theta) = \theta^\top K \theta$, where K is data covariance.
- $\kappa(H)$ controls gradient convergence speed for the probe.
- Therefore, shaping θ can reshape κ for downstream tasks.

- Singular values of $H(\theta)$ depend on alignment between θ and K 's singular vectors.
- Selective immunization is feasible when K_P and K_H are sufficiently misaligned.
- Angles between principal directions dictate how κ can be steered.
- Takeaway: geometry of data covariances governs achievable selectivity.

- Criteria: (a) $\kappa(H_H(\theta_I))$ increases vs $\kappa(H_H(I))$ on harmful tasks.
- (b) $\kappa(H_P(\theta_I))$ does not increase vs $\kappa(H_P(I))$ on pre-training.
- (c) Pre-training performance is maintained.
- Metric: Relative Immunization Ratio (RIR) $\gg 1$ indicates success.

- Objective: $\mathcal{R}_{\text{ill}}(H_H(\theta)) + \mathcal{R}_{\text{well}}(H_P(\theta)) + \mathcal{L}(\mathcal{D}_P; \theta, \omega)$.
- Increase κ for harmful task; decrease κ for pre-training task; keep DP loss low.
- Closed-form gradients through $H(\theta) = \theta^\top K \theta$ enable end-to-end training.
- Covariance-normalized updates preserve κ monotonicity under conditions.

Method visual: κ predicts convergence (illustration)

- They illustrate how larger κ slows gradient convergence and smaller κ accelerates it.
- The plotted norm-ratio dynamics are consistent with the Hessian condition-number predictions.
- Takeaway: steering κ via θ translates into predictable optimization behavior.




Figure: Norm ratio Eq. (16) vs. epochs under gradient descent with exact line search.

Component 1: κ -minimizing regularizer R_{well}

- Purpose: improve conditioning for the pre-training task.
- Property: monotonic κ decrease under gradient descent (from prior work).
- Effect: accelerates DP optimization and preserves utility.
- Used on $H_P(\theta)$ to protect benign adaptation speed.

- Form: $R_{\text{ill}}(S) = \frac{1}{\frac{1}{2}k^{-1}\|S\|_F^2 - \frac{1}{2}(\sigma_{\min}(S))^2}$.
- Nonnegative; upper-bounds $1/\log \kappa$; differentiable when $\sigma_{\min}$ is unique.
- Gradient steps provably increase κ for suitable step sizes.
- Applied to $H_{\text{H}}(\theta)$ to slow or destabilize harmful fine-tuning.

- Closed-form ∇_{θ} for $H(\theta) = \theta^{\top} K \theta$ (both regularizers).
- Covariance-normalized updates; mild uniqueness/stepsize assumptions.
- Dummy layer for linearizing deep encoders (ResNet18, ViT).

- Datasets: House Prices (regression), MNIST binary pairs (90), ImageNet pre-training vs Cars/Country211.
- Metrics: κ changes per task and Relative Immunization Ratio (RIR).
- Baselines: Rill Only (this work, ablation), IMMA (Zhou et al., NeurIPS 2023), Opt κ (this work, direct κ optimization), Init θ_0 (identity).
- Goal: achieve $\text{RIR} \gg 1$ while preserving DP accuracy.

- They report an average RIR ≈ 70.0 with success across all 90 pairs.
- Opt κ shows high variance and fails in many pairs; the method is consistent.
- The method preserves or improves benign conditioning while degrading harmful.
- Shows broad applicability across diverse task pairs.
- Baselines as in "Experimental setup"; IMMA: Zhou et al., NeurIPS 2023.

Table: Table 2. Quantitative results of immunization in MNIST (LeCun, 1998), computed over 3 random seeds and averaged over all digit pairs.

Method	Eq. (15) (i)↑	Eq. (15) (ii)↓	RIR↑
Rill Only	14.832 ± 1.039	8.654 ± 0.606	1.933 ± 0.046
IMMA	4.522 ± 0.139	2.774 ± 0.094	1.774 ± 0.041
Opt κ	3.196 ± 1.225	0.756 ± 1.171	69.73 ± 54.00
Ours	6.345 ± 0.188	0.149 ± 0.009	70.04 ± 3.280

- The method achieves $RIR_{\theta_0} \gg 1$ on both ResNet18 and ViT.
- ViT retains or improves ImageNet accuracy (utility preserved).
- ResNet18 shows modest DP accuracy impact with strong RIR gains.
- Confirms transfer of linear-theory insights to nonlinear deep nets.
- Baselines as in “Experimental setup”; IMMA: Zhou et al., NeurIPS 2023.

Table: Table 3. Quantitative results of immunization of model pre-trained on ImageNet (Deng et al., 2009), computed over 3 random seeds.

DH	Method	ResNet18				ViT			
		Eq. (17) (i)†	Eq. (17) (ii)‡	RIR_{θ_0} †	DP Test Acc. (%)†	Eq. (17) (i)†	Eq. (17) (ii)‡	RIR_{θ_0} †	DP Test Acc. (%)†
Cars	Intr. θ_0	1.0	1.0	1.0	68.24	1.0	1.0	1.0	81.78
	RIR Only	1.878 ± 0.034	1.796 ± 0.025	1.057 ± 0.026	63.84 ± 0.292	13.121 ± 0.038	4.097 ± 0.098	3.342 ± 0.046	82.31 ± 0.035
	IMMA	0.866 ± 0.002	0.899 ± 0.001	0.974 ± 0.002	63.57 ± 0.234	1.422 ± 0.006	2.090 ± 0.043	0.702 ± 0.007	81.89 ± 0.010
	Opt κ	1.217 ± 0.021	0.798 ± 0.005	1.527 ± 0.019	63.65 ± 0.148	3.598 ± 0.510	0.171 ± 0.033	26.369 ± 2.814	82.51 ± 0.085
	Ours	2.386 ± 0.442	0.699 ± 0.062	3.467 ± 0.358	62.36 ± 0.173	7.045 ± 0.247	0.323 ± 0.086	34.517 ± 0.896	82.79 ± 0.200
Country211	RIR Only	20.727 ± 0.791	20.675 ± 1.685	1.038 ± 0.050	82.17 ± 1.599	69.292 ± 1.198	63.519 ± 6.620	1.122 ± 0.007	80.73 ± 0.129
	IMMA	0.791 ± 0.005	0.814 ± 0.006	0.972 ± 0.007	87.03 ± 0.146	6.242 ± 0.203	7.598 ± 0.717	0.845 ± 0.046	82.47 ± 0.036
	Opt κ	1.538 ± 0.155	1.053 ± 0.091	1.472 ± 0.043	86.81 ± 0.115	4.589 ± 0.079	0.300 ± 0.106	16.498 ± 5.183	82.79 ± 0.023
	Ours	3.267 ± 0.330	0.399 ± 0.034	8.714 ± 0.672	85.01 ± 0.143	20.894 ± 1.425	0.700 ± 0.082	41.341 ± 0.967	83.37 ± 0.075

- Feasibility hinges on alignment angles between K_P and K_H .
- Stronger misalignment $\Rightarrow$ easier to increase DH κ while reducing DP κ .
- Limits: highly aligned tasks reduce achievable selectivity.
- In practice, the algorithm shows gains on deep models despite nonlinearity.

Limitations and practical considerations

- Uniqueness of extreme singular values is assumed for monotonicity proofs.
- Step-size tuning and covariance normalization are important in practice.
- Definition focuses on linear probes; nonlinear adapters may partially circumvent.
- Trade-offs: small DP accuracy drops are possible on some backbones.

Acknowledgments

- They thank collaborators, advisors, and supporting institutions.
- Community resources and open datasets enabled this research.

- IMMA: Zhou, Y., Qin, C., Duan, R., Raghunathan, A., and Kolter, J. Z. IMMA: Immunizing Text-to-Image Diffusion Models Against Malicious Fine-Tuning. Advances in Neural Information Processing Systems (NeurIPS), 2023.
- Other baselines are internal to this work (Rill Only, Opt κ , Init θ_0).

- How can the theory extend to nonlinear extractors?
- What scheduling balances $\kappa(H_P)$ and $\kappa(H_H)$ automatically?
- Can multi-harm settings be handled without sacrificing DP utility?
- How does this interact with concept erasure and tamper-resistance?