

Open-YOLO 3D: Towards Fast and Accurate Open-Vocabulary 3D Instance Segmentation

Mohamed El Amine Boudjoghra, Angela Dai, Jean Lahoud, Hisham Cholakkal, Rao Muhammad Anwer, Salman Khan, Fahad Shahbaz Khan

Mohamed Bin Zayed University of Artificial Intelligence (MBZUAI)
Technical University of Munich (TUM)
Aalto University
Australian National University
Linköping University

Background: 3D scene understanding is transforming the world

- 3D perception powers robotics, AR, and digital twins.
- Instance segmentation enables object-level reasoning and action.
- Open-vocabulary enables naming unseen objects from text prompts.
- Fast, accurate 3D perception is critical for real-time systems.

The problem: Open-vocabulary 3D instance segmentation at speed

- Task: segment 3D instances and label via text prompts.
- Recent solutions rely on SAM and CLIP feature lifting.
- Issue: multi-view 2D→3D aggregation is slow and costly.
- Consequence: unacceptable latency for real applications.

Key challenges in prior art

- Heavy 2D segmentation inflates compute per frame.
- CLIP feature lifting requires dense multi-view fusion.
- Per-instance visibility handled iteratively, not batched.
- Semantic-only scoring often up-ranks bad masks.

- Detector-only prompting: replace SAM+CLIP with OV detectors.
- Low-Granularity label maps: bounding-box fills, not masks.
- VAcc: vectorized, GPU-friendly visibility computation.
- MVPDist: multi-view voting for robust text labeling.
- Joint mask confidence: geometry $\times$ semantics without CLIP.

- 3D proposals, 2D detections, batched visibility, multi-view votes.
- Assign labels via MVPDist and score via $\text{sIoU} \cdot \text{sclass}$.
- No SAM masks, no CLIP feature lifting.

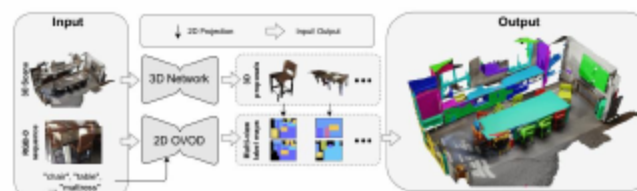


Figure: VAcc assigns top-k visible views; MVPDist aggregates labels at 2D projections to choose the class.

- Takeaway: VAcc + MVPDist turn multi-view loops into batched, GPU-efficient steps.

3D proposals: class-agnostic Mask3D

- Generate K class-agnostic 3D masks per scene.
- Apply NMS to filter overlapping proposals.
- Projections already encode instance structure.
- Avoids costly 2D-to-3D instance fusion.

2D detections and LG label maps

- Run YOLO-World on each RGB frame.
- Rasterize boxes to frame-sized LG label maps.
- Larger boxes overwrite to mimic occlusions.
- Cheap to store, fast to sample per projection.

- Project all points to all views in one GPU pass.
- Compute in-frame and depth-consistent visibility.
- Aggregate per-mask, per-view visibility scores.
- Select top-k views per mask efficiently.

- Sample LG labels at mask projections in top-k views.
- Aggregate class occurrences into a distribution.
- Assign class with maximum probability.
- Redundancy across views boosts zero-shot robustness.

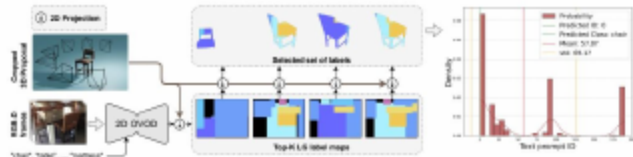


Figure: MVPDist samples labels from top-k visible views and selects the majority class.

- Takeaway: MVPDist picks the most frequent, view-consistent class across visible projections.

- Score: $sm = sloU \times sclass$.
- $sclass$: MVPDist's max class probability.
- $sloU$: projected 3D box vs best 2D box IoU across views.
- Downranks leaking or mislocalized masks.



Figure: $sloU$ = IoU of projected 3D box vs best 2D detector box across views; covers three typical cases.

- Takeaway: $sloU$ grounds semantics in geometry to avoid overconfident but spatially wrong masks.

- LG maps: boxes suffice; 3D masks carry geometry.
- HG via SAM increases latency drastically.
- LG matches HG predictions with less compute.
- Top-k aggregation stabilizes labels further.

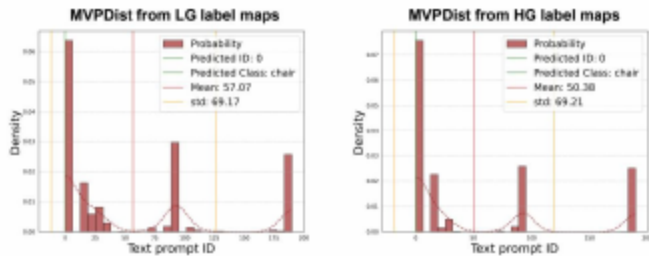


Figure: LG achieves similar predictions to HG with 5× lower runtime.

- Takeaway: LG label maps preserve accuracy while slashing runtime.

- ScanNet200: 312 val scenes, long-tail splits.
- Replica: 48 categories, cross-dataset generalization.
- Metrics: mAP@[0.5:0.95], mAP50, mAP25.
- Frames sampled as in OpenMask3D; single A100 GPU.

- Compare to Open3DIS, OpenMask3D, OpenScene, OVIR-3D.
- Proposal source: Mask3D for fairness and speed.
- Evaluate both accuracy and seconds per scene.
- Two scenarios: oracle masks, proposal masks.

- Open-YOLO 3D: mAP 24.7, mAP50 31.7, mAP25 36.2 at 21.8 s/scene.
- Beats prior CLIP+SAM pipelines in both AP and latency (up to 16× faster).
- Robust across head/common/tail: 27.8/24.3/21.6.

- Open-YOLO 3D: mAP 23.7, mAP50 28.6, mAP25 34.8 at 16.6 s/scene.
- About 11× faster than CLIP+SAM; strong cross-dataset generalization.

Runtime breakdown: where we gain speed

- Visibility and labeling dominate prior runtimes.
- VAcc collapses iterative loops into one pass.
- Detector-only LG maps remove SAM overhead.
- Overall: seconds per scene, not minutes.

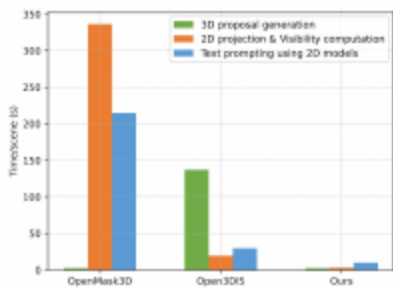


Figure: VAcc removes per-mask loops; detector-only LG maps avoid SAM overhead.

- Takeaway: Batched projection and depth filtering remove per-mask loops.

Replica qualitative results

- Tighter, cleaner instance masks with fewer spurious regions.
- Better adherence to object boundaries; higher precision.
- Consistent under varied lighting and layouts; strong generalization.

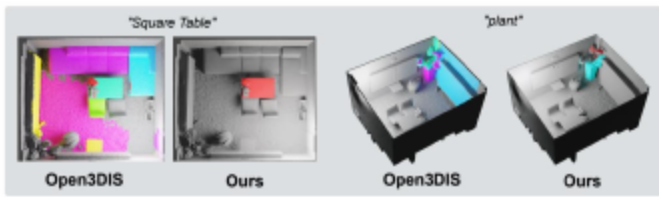


Figure: Joint geometry + semantics improves boundary adherence and precision.

- Takeaway: Joint geometry + semantics improves visual quality across diverse scenes.

Qualitative: open-vocabulary on ScanNet200

- Sharper masks than Open3DIS across prompts.
- Handles occlusions and small objects with improved robustness.
- Zero-shot labels align with text prompts.

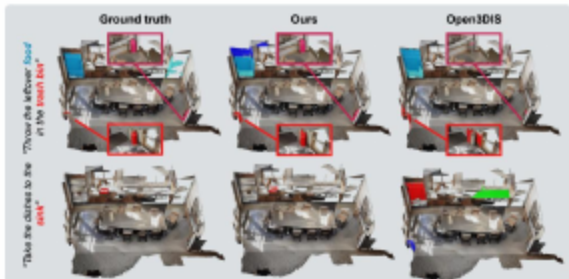


Figure: Multi-view voting stabilizes labels under occlusion and sparse views.

- Takeaway: Multi-view voting adds stability under occlusion and sparse views.

Scoring quality: geometry × semantics matters

- s_{class} = MVPDist's max class probability; s_{IoU} = 2D box IoU across views from the 3D mask projection.
- Effect: penalizes leakage/mislocalization; promotes well-localized masks.
- Semantic-only methods overvalue bad masks; our ranking elevates well-localized instances.



Figure: s_{IoU} - s_{class} aligns mask scores with human-perceived quality; semantic-only scores overvalue leaky masks.

- Takeaway: s_{IoU} - s_{class} aligns scores with human-perceived mask quality.

- Small-object recall limited by 3D proposals.
- Detector errors can bias MVPDist votes.
- Heavy occlusions reduce informative views.
- Vocabulary coverage depends on detector training.

- We thank collaborators, institutions, and funding agencies.
- Community projects: Mask3D, YOLO-World, and benchmarking datasets.

- How can we boost small-object recall without re-introducing heavy 2D segmentation?
- Can we train end-to-end across detection–projection–labeling for further gains?
- How do we expand vocabulary coverage under domain shift and long tails?
- What’s the path to true online/streaming OV 3D segmentation?
- Can lightweight 2D instance cues be fused for better occlusion handling?