

Rethinking Post-Hoc Search-Based Neural Approaches for Solving Large-Scale Traveling Salesman Problems

Yifan Xia, Xianliang Yang, Zichuan Liu, Zhihao Liu, Lei Song, Jiang Bian


2024

01 Introduction: Heatmap-Guided MCTS—Questions and Contributions.


- ML heatmaps + MCTS face train–test mismatch and scaling limits.
- SoftDist provides a supervision-free heatmap tuned with MCTS for **aligned objectives**.
- A Score metric reveals **persistent inefficiency** of MCTS versus **LKH-3** under equal compute.
- Experiments show SoftDist matches or surpasses most **ML heatmaps** with lower latency.
- Positioning: prior heatmaps train without MCTS; we address this misalignment.

02 SoftDist: Simple Distance-Based Heatmap and MCTS-Aligned Tuning

- Heatmap: softmax on distances with one **temperature** τ .
- Tune τ by grid search with **MCTS-in-the-loop** to minimize tour length.
- **Label-free**, fast, and training–inference **aligned**.



03 Preliminaries: TSP and How MCTS Consumes Heatmaps



- 2D Euclidean **TSP**: minimize tour length.
- Models output an $n \times n$ **heatmap** guiding choices.
- Within **MCTS**: construct tour, sample guided moves, apply k-opt.
- Backpropagate improvements; our focus is **inference-time** use of heatmaps.

04


Evaluation Protocol: Fair LKH-3 Head-to-Head, Score Metric, Datasets, and Settings

- A compute-equal head-to-head matches wall-clock for MCTS and LKH-3.
 - The **Score** equals the gap of LKH-3 divided by the gap of MCTS, with higher favoring MCTS.
 - Datasets** include TSP 500, TSP 1000, and TSP 10000 drawn from public benchmarks.
- Tuning:** Temperature values are tuned by grid search on held out sets with MCTS in the loop.
 - Metrics:** tour length, gap, total time, and the Score under matched budgets.

06

Sensitivity: MCTS Settings and Time Budgets

- Tuned** MCTS settings improve most methods yet do not close the gap to **LKH-3**.
- With more time, LKH-3 approaches optimal solutions and MCTS Scores **decrease**.
- On **TSP-10000**, SoftDist performs best among MCTS variants yet remains below **LKH-3**.
- Sensitivity indicates reduced **heatmap influence** under stronger MCTS tuning.
- Time budget curves **plateau** for large instances, showing limited MCTS returns.



METHOD	TSP-500				TSP-1000			
	LENGTH ↓	GAP ↓	TIME ↓	SCORE ↑	LENGTH ↓	GAP ↓	TIME ↓	SCORE ↑
LKH-3 (DEFAULT)	16.55*	0.00%	46.28m	—	23.12*	0.00%	2.57h	—
ATT-GCN [†]	16.72	1.02%	0.52m+0.67m	5.38%	23.48	1.58%	0.73m+1.43m	13.77%
DIMES [‡]	16.75	1.26%	0.97m+0.68m	4.35%	23.61	2.11%	2.08m+1.45m	10.29%
DIFUSCO [#]	16.69	0.90%	3.61m+0.68m	6.12%	23.48	1.56%	11.86m+1.43m	13.93%
UTSP [†]	16.73	1.09%	1.37m+0.68m	5.05%	23.50	1.65%	3.35m+1.45m	13.18%
SoftDist	16.72	1.03%	0.00m+0.68m	5.32%	23.52	1.73%	0.00m+1.44m	12.56%
ZEROS	16.72	1.06%	0.00m+0.68m	5.20%	23.55	1.85%	0.00m+1.44m	11.72%

05

Results: SoftDist vs ML Heatmaps; MCTS vs LKH-3 via Score

- SoftDist matches or surpasses most machine learned heatmaps with **minimal latency**.
- A diffusion based method is competitive but **computationally costly**, while SoftDist is **label-free**.
- The Score shows MCTS far below **LKH-3** under equal compute across datasets.
- On large instances, SoftDist remains Pareto competitive in quality versus latency.
- Reported gaps are small for SoftDist but remain above **LKH-3**.
- SoftDist gaps: **1.44%/2.20%/3.13%** (500/1000/10000).

07

Conclusion and Outlook: Rethinking Heatmap-Guided TSP Solving

- SoftDist is competitive in quality and fast, while MCTS trails **LKH-3** by the Score metric.
 - Misaligned surrogate training **undermines practicality** of heatmap guided search.
 - Future directions include theory grounded objectives and **end to end** solvers.
- The evidence questions the practical value of **post hoc** neural search for large scale TSP.
 - We advocate simpler baselines and fair compute matched comparisons in future studies.

Thank You

Questions & Discussion