

Position: Rethinking Post-Hoc Search-Based Neural Approaches for Solving Large-Scale Traveling Salesman Problems

Yifan Xia, Xianliang Yang, Zichuan Liu, Zhihao Liu, Lei Song, Jiang Bian

Nanjing University, Nanjing, China
Microsoft Research Asia, Beijing, China
Institute of Automation, Chinese Academy of Sciences, Beijing, China

Combinatorial optimization already impacts real systems at scale

- Routing, logistics, and network design drive global efficiency and sustainability.
- The Traveling Salesman Problem (TSP) is a cornerstone with real-world impact.
- ML is increasingly combined with operations research to tackle scale.

Y. Xia et al. (Nanjing University) Position: Rethinking Post-Hoc Neural TSP... 1 / 20

Y. Xia et al. (Nanjing University) Position: Rethinking Post-Hoc Neural TSP... 2 / 20

Large-scale TSP demands reliable solutions under tight compute budgets

- Goal: high-quality tours for 500–10,000 nodes under realistic resources.
- Trend: heatmap-guided MCTS uses ML-predicted edge priors to refine tours.
- Question: does this guidance add value over strong heuristics?




Figure: Heatmap-guided MCTS phases.

Surrogate objectives and unfair budgets undermine heatmap-guided MCTS

- Training losses often misalign with the search used at inference.
- Resource fairness is unclear versus strong OR baselines like LKH-3.
- Uniform or near-zero heatmaps can perform similarly in practice.

Y. Xia et al. (Nanjing University) Position: Rethinking Post-Hoc Neural TSP... 3 / 20

Y. Xia et al. (Nanjing University) Position: Rethinking Post-Hoc Neural TSP... 4 / 20

A simple, aligned baseline and fair scoring overturn prior claims

- Provide a systematic theoretical and empirical critique of heatmap-guided MCTS.
- Introduce SoftDist: a supervision-free, objective-aligned heatmap with negligible latency.
- Define a resource-fair comparison protocol and gap-ratio score to benchmark against LKH-3.
- Show that under matched budgets, LKH-3 remains strongest; SoftDist matches prior ML quality at far lower cost.

A one-parameter distance-based heatmap rivals complex ML models

- Compute $\Phi_{ij} = \text{softmax}_j(-d_{ij}/\tau)$ so shorter edges get higher probability.
- A single temperature controls exploration versus greediness; no labels and negligible cost.
- In practice, this minimal prior competes with learned heatmaps at a fraction of the latency.

MCTS uses heatmaps only as weak priors for local k-opt search

- An ML heatmap provides edge-inclusion priors; MCTS biases local improvements accordingly.
- Default MCTS parameters are shared across heatmap sources for fairness and a balanced search.
- Search gains are backpropagated in the tree; the heatmap influences but does not dictate moves.

Optimizing τ on the actual search loop closes the train-test gap

- Select the temperature per scale by grid-searching directly on the same MCTS used at test time.
- This aligns the objective with the deployed search, avoiding training-inference mismatch.
- A short, coarse grid search finds near-optimal settings per scale; sensitivity is mild except at very small temperatures on very large instances.




Figure: SoftDist-guided MCTS phases with temperature tuning in the loop.

Matched hardware and a gap-ratio metric enable fair comparisons

- Gap versus a strong baseline is the relative excess tour length.
- Score compares MCTS and LKH-3 by the ratio of their gaps; higher means MCTS is closer to the baseline.
- Setup: synthetic Euclidean TSP-500/1000/10000; shared MCTS configs and CPU; SoftDist generation on GPU adds negligible latency.
- Interpretation: 100% means parity; below 100% favors LKH-3; above 100% favors MCTS.

SoftDist matches quality at negligible latency

- X-axis sweeps heatmap generation time; SoftDist sits near zero while maintaining strong tour quality.
- Headline: comparable or better quality than prior ML heatmaps with orders-of-magnitude lower latency.




Figure: Lower inference latency (x) and lower solution cost (y) are better; SoftDist is near-zero latency with competitive quality.

Across scales, SoftDist is competitive; LKH-3 remains strongest

- SoftDist+MCTS achieves small gaps across 500/1000/10000 nodes with negligible heatmap time.
- DIFUSCO attains similar quality but requires heavy supervision and long inference time.
- Strong OR (LKH-3) is still the most competitive in both quality and efficiency.

Table: Table 1. Results on large-scale TSP problems.

Method	Type	TSP-500			TSP-1000			TSP-10000		
		Length ↓	Gap ↓	Time ↓	Length ↓	Gap ↓	Time ↓	Length ↓	Gap ↓	Time ↓
CONCORDE	OR(EXACT)	16.55*	—	37.66M	23.12*	—	6.89H	N/A	N/A	N/A
GUROBI	OR(EXACT)	16.55	0.00%	45.03H	N/A	N/A	N/A	N/A	N/A	N/A
LKH-3 (DEFAULT)	OR	16.55	0.00%	46.28M	23.12	0.00%	2.57H	71.78*	N/A	N/A
LKH-3 (LESS TRAILS)	OR	16.55	0.00%	3.03M	23.12	0.00%	7.73M	71.79	—	51.27M
NEAREST INSERTION	OR	20.62	24.59%	05	28.98	25.26%	05	80.51	26.11%	05
RANDOM INSERTION	OR	18.57	12.21%	05	26.12	12.96%	05	81.85	14.04%	45
FARTHEST INSERTION	OR	18.30	10.57%	05	25.72	11.25%	05	80.59	12.29%	65
EAN	RL+S	28.63	73.03%	20.18M	50.30	117.59%	37.07M	N/A	N/A	N/A
EAN	RL+S+2-OPT	23.75	43.57%	57.70M	47.73	109.46%	5.30H	N/A	N/A	N/A
AM	RL+S	22.64	36.84%	15.64M	42.80	85.15%	83.97M	431.58	500.27%	12.63M
AM	RL+G	20.02	20.99%	1.51M	31.15	34.75%	3.10M	141.08	97.39%	5.99M
AM	RL+BS	19.53	18.03%	21.99M	29.90	29.23%	1.64H	129.40	80.28%	1.81H
GCN	SL+G	26.72	79.81%	0.67M	48.62	110.29%	26.52M	N/A	N/A	N/A
GCN	SL+BS	30.37	83.55%	38.02M	51.26	121.73%	51.87M	N/A	N/A	N/A
POMO+EAS-EMB	RL+AS	19.24	18.25%	12.80H	N/A	N/A	N/A	N/A	N/A	N/A
POMO+EAS-LAY	RL+AS	19.15	18.92%	16.19H	N/A	N/A	N/A	N/A	N/A	N/A
POMO+EAS-TAB	RL+AS	24.54	48.22%	11.63H	49.58	114.36%	63.45H	N/A	N/A	N/A
DIFUSCOz	SL+MCTS	16.63	0.51%	3.83M+1.67M	23.39	1.18%	11.86M+3.34M	73.76	2.77%	28.51M+16.87M
ATT-GCN	SL+MCTS	16.82	1.64%	0.52M+1.67M	23.67	2.37%	0.73M+3.34M	74.50	3.80%	4.36M+16.77M
DIMES	RL+MCTS	16.84	1.77%	0.52M+1.67M	23.68	2.40%	0.73M+3.34M	74.50	3.81%	4.36M+16.77M

Even with longer budgets, MCTS trails LKH-3 across scales

- X-axis sweeps wall-clock time; LKH-3 improves faster and sustains a lead.
- Headline: longer budgets do not close the gap; SoftDist is best among MCTS but still behind LKH-3.




Figure: Performance of MCTS under default settings for TSP-500, TSP-1000, and TSP-10000.

Tuning MCTS barely changes the story across heatmaps

- X-axis sweeps time under alternative parameters; relative ordering among heatmaps changes little.
- Headline: SoftDist remains competitive without latency; DIFUSCO is strong but slow; a zero heatmap is close to learned heatmaps.




Figure: Under alternative MCTS parameters (Min et al., 2023), relative ordering across heatmaps barely changes.

Under matched resources, MCTS efficiency lags far behind LKH-3

- Gap-ratio Scores are far below LKH-3 on 500/1000 nodes, and remain low at 10k nodes.
- More supervision and heavier GPUs do not guarantee practical wins.

Table: Table 2. Resource consumption and Score comparison of various methods on TSPs.

Method	Supervision	Hardware	TSP-500	TSP-1000	TSP-10000
ATT-GCN†	✓	GTX 1080 TI GPU	0.74%	3.87%	24.66%
DIMES†	×	NVIDIA P100 GPU	0.68%	3.75%	28.99%
UTSP†	×	NVIDIA V100 GPU	0.35%	2.08%	—
DIFUSCO#	✓	8×NVIDIA V100 GPUS	2.39%	7.78%	33.82%
SOFTDIST	×	NVIDIA A100 GPU	0.84%	4.17%	29.88%

Near-zero heatmaps rival learned ones—heatmap quality has limited leverage

- Uniform and near-zero priors can rival learned heatmaps under current MCTS designs.
- On small scales, giving more time mainly benefits LKH-3, shrinking relative gains from heatmaps.
- Overall leverage of heatmap quality is modest compared with search design and budget.

Findings may shift beyond Euclidean TSP and our implementation choices

- Results are on synthetic Euclidean instances; other metrics or data may differ.
- Hardware and engineering details influence wall-clock comparisons.
- Scope is a position with baselines and critiques, not a new end-to-end solver.

Theory-aligned objectives and end-to-end ML offer the path forward

- Design objectives that align heatmap learning with search-time performance.
- Close training–inference gaps using differentiable or faithful search proxies.
- Pursue autonomous, generalizable, end-to-end solvers beyond post-hoc search.

Open questions: objectives, fairness, and generalization

- What objectives best align learning with the deployed search?
- How should we enforce resource fairness across diverse hardware stacks?
- Can end-to-end ML solvers generalize beyond Euclidean assumptions?

For practice today, choose LKH-3; if using MCTS, default to SoftDist

- For large-scale TSP under matched resources, prefer LKH-3 for reliability and speed.
- If using heatmap-guided MCTS, SoftDist is a strong default with negligible latency.
- Always weigh heatmap generation time against search time in end-to-end budgets.

Acknowledgments

- We thank collaborators and the open-source community.
- Code available: https://github.com/xyffffff/rethink_mcts_for_tsp
- Appreciation for prior OR and ML contributions that enabled this study.