

Rethinking Image-Scaling Attacks: The Interplay Between Vulnerabilities in Machine Learning Systems

Yue Gao Ilia Shumailov Kassem Fawaz

University of Wisconsin–Madison, Madison, WI, USA

Vector Institute, Toronto, ON, Canada

May 24, 2026

Navigation icons

Yue Gao, Ilia Shumailov, Kassem Fawaz (UW–Madison)

Rethinking Image-Scaling Attacks...

May 24, 2026

1 / 25

Robustness must hold across the entire vision pipeline

- Vision AI powers phones, cars, healthcare, retail, and safety.
- Deployed models process diverse, real-world images at scale.
- Robustness must hold across full data pipelines, not just models.

Navigation icons

Yue Gao, Ilia Shumailov, Kassem Fawaz (UW–Madison)

Rethinking Image-Scaling Attacks...

May 24, 2026

2 / 25

Every deployed model relies on upstream image scaling

- Deep models require fixed-size inputs (e.g., 224×224).
- Upstream scaling maps high-resolution to model size.
- Scaling reweights pixels; kernels vary by algorithm.

Scaling and model vulnerabilities interact end-to-end

- Two families: adversarial examples and scaling attacks.
- Independent defenses miss joint, end-to-end effects.
- Need to evaluate the entire pipeline together.




Figure: Illustration: an HR image that looks benign becomes a target image after scaling.

Navigation icons

Yue Gao, Ilia Shumailov, Kassem Fawaz (UW–Madison)

Rethinking Image-Scaling Attacks...

May 24, 2026

3 / 25

Navigation icons

Yue Gao, Ilia Shumailov, Kassem Fawaz (UW–Madison)

Rethinking Image-Scaling Attacks...

May 24, 2026

4 / 25

Decision-only APIs force query-efficient, gradient-free attacks

- Commercial systems return only labels, no scores.
- No gradients; query budgets are tight.
- Pre-scaling defenses further distort attack signals.

Attacking the full pipeline uncovers hidden, post-scaling perturbations

- Exploit scaling to hide perturbations post-resize.
- Guide black-box search using scaling-aware noise.
- Bypass pre-scaling defenses with defense-aware tricks.



Figure: Scaling-aware noise sampling guides black-box attacks to exploit scaler weaknesses.

Scaling-aware, defense-aware methods strengthen black-box attacks

- Scaling-aware Noise Sampling (SNS) for decision-based attacks.
- Defense-aware methods: median gradients and offline EOT.
- System-level analysis across models, defenses, and an API.
- Efficient SNS matches precise inversion with lower compute.

Optimizing in HR for post-scaling effects improves attack efficacy

- Pipeline: $F = f \circ g \circ h$ (defense, scaling, model).
- Operate in HR space but optimize for post-scaling effect.
- Plug-and-play with HSJ and Sign-OPT attacks.

Sampling in LR and projecting to HR targets the scaling-hidden subspace

- Sample noise in LR space where the model operates.
- Project LR noise back to HR via scaling pipeline.
- Bias search toward scaling-hidden directions.




Figure: SNS samples in LR and projects to HR to target scaling-hidden directions.

Efficient SNS projection guides HSJ and Sign-OPT

- Use projection-matching gradients as an efficient inverse; preserve query budget.
- HR HSJ: SNS informs boundary-direction estimates; HR Sign-OPT: optimize along SNS-guided directions.
- Maintain minimal post-scaling perturbations with no algorithm changes.

Approximating the median restores usable gradients

- Median flattens response, stalling line search.
- Remedy: trimmed, weighted average approximation.
- Model as masked pooling to target vulnerable pixels.

Offline EOT captures randomized filtering with zero extra queries

- Direct EOT is query-heavy in black-box.
- Fold randomness into expected scaling offline.
- Zero extra queries; uniform weighting in expectation.

Joint optimization outperforms sequential pipelines

- Naive: craft LR adversarial, then hide via scaling.
- Joint: optimize end-to-end with SNS guidance.
- Joint approach wins in success and perturbation.




Figure: Joint optimization outperforms a sequential “attack then scale” pipeline.

ImageNet/CelebA and ResNets provide a realistic evaluation

- Datasets: ImageNet, CelebA.
- Models: ResNet-50/34; inputs resized to 224×224.
- Metrics: scaled ℓ_2 , attack success rate (ASR), queries.

HR variants achieve lower scaled ℓ_2 and higher ASR without defenses

- Lower scaled ℓ_2 at same query budgets.
- Higher ASR for HR HSJ and HR Sign-OPT.
- Ablation: removing SNS erases the gains.




Figure: HR variants achieve lower scaled ℓ_2 and higher ASR at the same queries.

Median filtering does not stop SNS-based HR attacks

- HR retains strong gains under median filtering.
- Improved median gradients restore convergence.
- LR baselines degrade substantially in comparison.




Figure: SNS-based HR attacks retain gains even under median filtering.

HR HSJ halves ℓ_2 and roughly doubles ASR on a commercial API

- HR HSJ beats LR HSJ under same query budgets.
- Lower perturbation and higher success rates.
- Scaling vulnerabilities exist in practical systems.




Figure: HR HSJ reduces ℓ_2 and doubles ASR on a commercial API.

HR C&W transfers better to the API without confidence tuning

- HR C&W transfers more effectively to the API.
- Reported scores confirm genuine gains.
- No reliance on inflated confidence settings.




Figure: HR C&W transfers more effectively without tuning confidence.

Higher HR scaling ratios cut scaled ℓ_2 and boost ASR across budgets

- Bigger HR inputs yield lower scaled ℓ_2 .
- ASR rises sharply with more queries.
- HR dominates LR across budgets and sizes.

Attacks	Query = 10K		Query = 15K		Query = 20K	
	ℓ_2	ASR	ℓ_2	ASR	ℓ_2	ASR
HSJ (LR, 224×224)	5.26	10.8%	4.16	14.0%	3.58	18.3%
HSJ (HR, 672×672)	1.73	59.6%	1.37	74.5%	1.17	86.2%
HSJ (HR, 1120×1120)	1.03	90.4%	0.82	98.9%	0.70	100.0%

On the cloud API, HR HSJ achieves lower ℓ_2 and higher ASR at equal queries

- HR achieves lower ℓ_2 at same query budgets.
- ASR roughly doubles versus LR at 2K–3K queries.
- Confirms scalability to commercial services.

Attacks	Query = 1K		Query = 2K		Query = 3K	
	ℓ_2	ASR	ℓ_2	ASR	ℓ_2	ASR
HSJ (LR)	10.57	10.9%	6.64	18.2%	7.42	21.8%
HSJ (HR)	4.24	25.9%	2.93	43.1%	1.92	51.7%

White-box results: robust-by-design scaling resists, others break

- HR PGD/C&W confirm lower end-to-end perturbations across typical scalers.
- Median filtering can be bypassed with proper gradient approximations.
- Robust-by-design or randomized scaling neutralizes HR advantages.




Figure: Robust-by-design scaling removes the exploitable subspace; HR brings no extra gains.

Robustness is a pipeline property, not a component attribute

- End-to-end HR attacks exploit interactions between scaling and the model.
- Component-only defenses overestimate robustness; SNS generalizes across attacks and datasets.
- Only randomized or robust-by-design scaling is effective in practice.

Pipeline metrics, broader norms, and detectors remain open

- Pipeline-level robustness metrics and certification.
- Extend SNS to other black-box families, norms.
- Learning-based detectors for subtle, hidden perturbations.

Questions

- Thank you for your attention!
- Questions and feedback are welcome.

- We thank collaborators, advisors, and supporting institutions.
- We acknowledge open-source libraries and tools used.