

To Smooth or Not? When Label Smoothing Meets Noisy Labels

Jiaheng Wei, Hangyu Liu, Tongliang Liu, Gang Niu, Masashi Sugiyama, Yang Liu

University of California, Santa Cruz
Brown University
TML Lab, Sydney AI Centre, The University of Sydney
RIKEN AIP
University of Tokyo

Navigation icons

Jiaheng Wei et al. (UC Santa Cruz)

To Smooth or Not?...

1 / 24

Why we care: Machine learning with imperfect labels is transforming the world

- We scale modern ML by training on massive, cheaply labeled datasets (web, crowdsourcing).
- We routinely face imperfect labels in vision, NLP, and recommendation systems.
- We observe that noisy labels hurt generalization, slow training, and can miscalibrate confidence.

Navigation icons

Jiaheng Wei et al. (UC Santa Cruz)

To Smooth or Not?...

2 / 24

Our background: Label Smoothing (LS) as a ubiquitous regularizer

- We use LS [Szegedy et al., 2016] to replace one-hot labels with a mix of hard labels and the uniform distribution.
- We find benefits on clean data: better generalization and calibration across tasks.
- We ask: how should we regularize labels when they are already noisy?

Navigation icons

Jiaheng Wei et al. (UC Santa Cruz)

To Smooth or Not?...

3 / 24

Our question: When does LS help or hurt under noisy labels?

- We see reports that LS can improve robustness at low/moderate noise.
- We also see that at high noise rates, LS can oversmooth and wash out useful signal.
- We aim to give a principled rule to choose between LS, vanilla CE, or alternatives.

Navigation icons

Jiaheng Wei et al. (UC Santa Cruz)

To Smooth or Not?...

4 / 24

Our challenge: We need to shape confidence under noise

- We recognize that noisy labels reduce the effective confidence signal during learning.
- We note that extra smoothing can further depress confidence, leading to underfitting.
- We seek a mechanism that boosts confidence for correct patterns while suppressing it for wrong ones.

Our idea and contributions

- We unify positive LS and Negative Label Smoothing (NLS) via Generalized Label Smoothing (GLS).
- We show a theory-driven phase transition: LS at low noise; NLS at high noise.
- We derive a closed-form optimal smoothing r_{opt} and connect NLS to robust losses.
- We provide extensive evidence and a practical recipe (CE warm-up $\rightarrow$ NLS).

Our method: Generalized Label Smoothing (GLS)

- We define GLS: $y^{\text{GLS},r} = (1-r) \cdot y + (r/K) \cdot \mathbf{1}$, extended to $r \in (-\infty, 1]$.
- We recover LS when $r > 0$ (probability mass toward uniform) and vanilla when $r = 0$.
- We recover NLS when $r < 0$ (negative weighting; labels sum to 1 though entries may be negative).

How we choose LS vs NLS

- We choose r based on the noise rate e and the clean-data preference r^* .
- Our decision map (figure) yields LS when $r_{\text{opt}} > 0$, vanilla when $r_{\text{opt}} = 0$, and NLS when $r_{\text{opt}} < 0$.
- Reading guide: above the boundary is the NLS region; below is LS; the switch occurs where $r_{\text{opt}} = 0$.
- We observe on UCI datasets the same phase transition: as noise grows, the best r turns negative, matching r_{opt} predictions.




Figure: Our decision between NLS and LS given e and r^* .

Our lens: Risk equivalence under noisy labels (Theorem 3.2)

- We show that minimizing a loss with smoothed noisy labels equals the clean target risk plus bias terms.
- We analyze how the bias terms directly modulate model confidence under symmetric/asymmetric noise.
- We use this lens to explain why LS reduces and NLS increases confidence in targeted ways.

How we connect to robust learning methods

- We show loss correction (forward/backward) [Patrini et al., 2017] is recoverable as NLS with specific negative r .
- We link complementary labels [Ishida et al., 2017; Ishida et al., 2019] to the limit $r \rightarrow -\infty$ (an NLS view).
- We interpret Peer Loss [Liu and Tao, 2020] as an NLS-regularized objective under certain priors.
- We unify why these methods work: all shape confidence under noise.

Our phase transition formula for symmetric noise

- Under symmetric noise rate e , we obtain $r_{\text{opt}} = \frac{r^* - 2e}{1 - 2e}$.
- We see three regimes: $e < r^*/2 \Rightarrow$ LS ($r_{\text{opt}} > 0$); $e = r^*/2 \Rightarrow$ Vanilla; $e > r^*/2 \Rightarrow$ NLS ($r_{\text{opt}} < 0$).
- We note the multiclass extension follows similarly with ϵ .




Figure: Optimal smooth rates on UCI datasets vs. label noise (ties possible).

How we train: a practical recipe for NLS

- We warm up with CE to learn clean patterns, then switch to NLS.
- From ablations, direct NLS can collapse early (overconfidence on noise); warm-up stabilizes training and unlocks NLS gains.
- We find negative r robust at high noise; moderately negative values (e.g., $r \approx -0.4$ to -4) work well; larger $|r|$ may require longer warm-up.




Figure: Bias and variance of LS/Vanilla/NLS models on clean CIFAR-10.

- We evaluate on synthetic 2D, UCI (Image, Waveform, Heart, Twonorm, Banana), CIFAR-10/100, AGNews, CIFAR-N, Clothing1M.
- We vary noise types (symmetric/asymmetric) and noise rates.
- We measure test accuracy and analyze phase transitions and confidence calibration.
- Our baseline suite includes CE, LS [Szegedy et al., 2016], NLS (ours), loss correction [Patrini et al., 2017], SCE [Wang et al., 2019], APL [Ma et al., 2020], Peer Loss [Liu and Tao, 2020], ELR [Liu et al., 2020], Bootstrap [Reed et al., 2015], AUM [Pleiss et al., 2020], and an f -divergence method [Zhang and Sabuncu, 2018].

- Label Smoothing (LS): [Szegedy et al., 2016]
- Forward/Backward Loss Correction: [Patrini et al., 2017]
- Complementary Labels: [Ishida et al., 2017; Ishida et al., 2019]
- Peer Loss: [Liu and Tao, 2020]
- Symmetric Cross-Entropy (SCE): [Wang et al., 2019]
- Active Passive Loss (APL): [Ma et al., 2020]
- Early-Learning Regularization (ELR): [Liu et al., 2020]
- Bootstrap: [Reed et al., 2015]
- AUM: [Pleiss et al., 2020]
- f -divergence method: [Zhang and Sabuncu, 2018]

What we see under noise: confidence becomes the bottleneck

- We see that with high noise, LS becomes uncertain over large regions (undershoots).
- We see NLS sharpen confidence on correct regions and suppress it on wrong ones.
- We find this matches the confidence-bias interpretation from our theory.
- What to look for: left=NLS retains confident regions on the true manifold; right=LS diffuses confidence and blurs boundaries; vanilla is in between.




Figure: Confidence maps on synthetic Type 1 data at $\epsilon=0.4$ (left: NLS, middle: Vanilla, right: LS).

What we observe on CIFAR-10/100: NLS wins at higher noise

- We find that on CIFAR-10 (asymmetric), NLS is competitive and yields significant gains at $\epsilon = 0.2$.
- On CIFAR-100 (symmetric), at $\epsilon = 0.6$ NLS beats the best LS by ~ 5 points (very small p -values).
- We also see the same LS→NLS switch on AGNews as noise rises (supports generality beyond vision).

Table: Table 4. Test accuracy (mean $\pm$ std) comparisons on asymmetric noisy CIFAR-10, symmetric CIFAR-100 datasets

Smooth Rate	CIFAR-10 Asymmetric		CIFAR-100 Symmetric	
	$\epsilon = 0.2$	$\epsilon = 0.3$	$\epsilon = 0.4$	$\epsilon = 0.6$
$r = 0.8$	90.45 $\pm$ 0.06	87.83 $\pm$ 0.13	54.04 $\pm$ 0.93	39.50 $\pm$ 0.18
$r = 0.6$	90.41 $\pm$ 0.09	87.83 $\pm$ 0.13	52.72 $\pm$ 0.15	40.49 $\pm$ 0.07
$r = 0.4$	90.49 $\pm$ 0.10	87.90 $\pm$ 0.13	54.26 $\pm$ 0.07	41.57 $\pm$ 0.05
$r = 0.0$	88.32 $\pm$ 0.24	86.27 $\pm$ 0.32	48.03 $\pm$ 0.29	38.11 $\pm$ 0.14
$r = -0.4$	87.27 $\pm$ 1.83	88.33 $\pm$ 0.06	56.87 $\pm$ 0.08	43.70 $\pm$ 0.16

What we observe in the wild: Clothing1M and CIFAR-N

- We find that NLS beats LS and strong baselines on Clothing1M (+0.8% over LS).
- We observe that across CIFAR-N splits, NLS achieves best or tied-best results.
- We highlight that GLS is simple, drop-in, and effective beyond synthetic noise.

Note: F-div in the table refers to: [Zhang and Sabuncu, 2018].

Table: Table 6. Performance comparisons on Clothing 1M and CIFAR-N: results of baselines are obtained through the public leader-board.

Method	Clothing 1M	CIFAR-10N Aggre	CIFAR-10N Rand1	CIFAR-10N Worse	CIFAR-100N Fine
CE	68.94	87.77	85.02	77.69	55.50
BLC	69.13	88.13	87.14	77.61	57.14
FLC	69.84	88.24	86.88	79.79	57.01
PL	72.60	90.75	89.06	82.53	57.59
F-div	73.09	91.64	89.70	82.53	57.10
LS (best)	73.44	91.57	89.80	82.76	55.84
NLS (best)	74.24	91.97	90.29	82.99	58.59

How we affect representations: larger margins with NLS

- We observe that NLS increases inter-class separation in pre-logits.
- We see sharper confidence translate to cleaner class clusters.
- We take this as evidence that NLS restores discriminative structure under noise.
- What to look for: tighter, well-separated clusters with NLS in pre-logit space indicate restored structure.




Figure: Pre-logit embeddings on CIFAR-10 (left: NLS, middle: Vanilla, right: LS).

How we shape confidence on CIFAR-10

- We see NLS concentrate high confidence on correct predictions.
- We see it reduce confidence on wrong predictions, improving reliability.
- We note that LS distributes confidence more moderately, which can underfit under noise.
- What to look for: we see higher confidence for correct and lower for wrong with NLS, widening separation compared to LS.




Figure: Confidence distributions on CIFAR-10 test (NLS, Vanilla, LS; trained with symmetric 0.2 label noise).

Our limitations and practical considerations

- Estimating noise rate or r^* is nontrivial; very large negative r without warm-up can destabilize training.
- We note that state-of-the-art pipelines (e.g., ELR [Liu et al., 2020]) may still outperform GLS alone in some settings.
- We treat the r schedule as a hyperparameter search in practice.

- We unify LS and NLS in GLS and explain their effects via confidence shaping.
- We recommend a phase-aware strategy: use LS at low noise; switch to NLS as noise grows.
- We advocate a simple, effective recipe: CE warm-up $\rightarrow$ NLS with moderately negative r .
- We clarify why several robust methods work (as NLS instances).

- We aim for automatic r scheduling via noise estimation or validation signals.
- We plan to combine NLS with advanced noisy-label pipelines (selection, semi-supervised).
- We will extend to detection, segmentation, and language modeling with varied label quality.
- We encourage rethinking label regularization beyond the probability simplex.

- Thank you for your attention!
- We welcome questions and feedback.

- We thank collaborators, reviewers, and open-source contributors.
- We are grateful to dataset maintainers and the research community.